



Data Mining

Lecture 3 Preprocessing Methods



<http://www.informatik.uni-hamburg.de/WTM/>

Data Preprocessing

- Data Preprocessing: An Overview 
 - Data Quality
 - Major Tasks in Data Preprocessing
- Data Cleaning
- Data Integration
- Data Reduction
- Data Transformation
- Data Discretization

Data Quality: why preprocess the Data?

Measures for data quality: A multidimensional view

- Accuracy: correct or wrong, accurate or not
- Completeness: not recorded, unavailable, ...
- Consistency: some modified but some not, dangling, ...
- Timeliness: timely update?
- Believability: how trustable are the data are?
- Interpretability: how easily the data can be understood?

Why We should Clean Dirty Data



You DO have dirty data...



-Proprietary and Confidential-
Copyright © 2007 INTRICITY, LLC
All Rights Reserved

Major Tasks in Data Preprocessing

- ***Data cleaning***

- Fill in missing values, smooth noisy data, identify or remove outliers, and resolve inconsistencies

- ***Data integration***

- Integration of multiple databases, data cubes, or files


- ***Data reduction***

- Dimensionality reduction
- Data compression

- ***Data transformation*** and ***data discretization***

- Normalization
- Concept hierarchy generation

Data Preprocessing

- Data Preprocessing: An Overview
 - Data Quality
 - Major Tasks in Data Preprocessing
- Data Cleaning 
- Data Integration
- Data Reduction
- Data Transformation
- Data Discretization

Data Cleaning

- Data in the real world is “dirty” or incorrect, e.g., instrument faulty, human or computer error, transmission error
- **Incomplete**: lacking attribute values, lacking certain attributes of interest, or only aggregate data available
 - e.g., *Occupation*=“ ” (missing data)
- **Noisy**: containing noise, errors, or outliers
 - e.g., *Salary*=“-10” (an error)
- **Inconsistent**: containing discrepancies in codes or names, e.g.,
 - *Age*=“42”, *Birthday*=“03/07/2012”
 - Was rating “1, 2, 3”, now rating “A, B, C”
 - discrepancy between duplicate records
- **Intentionally imprecise** (e.g., *disguised missing data*)
 - Jan. 1 as everyone’s birthday?

Incomplete (Missing) Data

- Data is not always available
 - E.g., many tuples have no recorded value for several attributes, such as customer income in sales data
- Missing data may be due to
 - equipment malfunction
 - inconsistent with other recorded data and thus deleted
 - data not entered due to misunderstanding
 - certain data may not be considered important at the time of entry
 - not register history or changes of the data
- Missing data may need to be inferred

How to handle missing Data?

- Ignore the tuple: usually done when class label is missing (when doing classification) — not effective when the % of missing values per attribute varies considerably
- Fill in the missing value manually: tedious + infeasible?
- Fill it in automatically with
 - a global constant : e.g., “unknown”, a new class?!
 - the attribute mean
 - the attribute mean for all samples belonging to the same class: smarter
 - the most probable value: inference-based such as Bayesian formula or decision tree

Missing Data

- One possible interpretation of missing values – **“don’t care”** values:


$$X = \{1, ?, 3\}$$

→ for the second feature the domain is $[0, 1, 2, 3, 4]$:

$$\begin{aligned} X1 &= \{1, 0, 3\}, & X2 &= \{1, 1, 3\}, & X3 &= \{1, 2, 3\}, \\ X4 &= \{1, 3, 3\}, & X5 &= \{1, 4, 3\} \end{aligned}$$

- *Data miner* can generate model of **correlation between features**.
 - Different techniques possible: regression, Bayesian formalism, clustering, or decision tree induction.

Missing Data Replacement with Regression Analysis





- In general, replacement of missing values using a simple, artificial schema of data preparation is *speculative and often misleading*.
- It is best to generate multiple solutions of data mining **with and without features** that have missing values, and then make comparison, analysis and interpretation.

Noisy Data

- Noise: random error or variance in a measured variable
- Incorrect attribute values may be due to
 - technology limitation
 - faulty data collection instruments
 - data entry problems
 - data transmission problems
 - inconsistency in naming convention
- Other data problems which require data cleaning
 - duplicate records
 - incomplete data
 - inconsistent data

How to handle noisy Data?


- Binning
 - first sort data and partition into (equal-frequency) bins
 - then one can smooth by bin means, smooth by bin median, smooth by bin boundaries, etc.
- Regression
 - smooth by fitting the data into regression functions
- Clustering
 - detect and remove outliers
- Combined computer and human inspection
 - detect suspicious values and check by human (e.g., deal with possible outliers)




Data Quality: why preprocess the Data?



Data Preprocessing

- Data Preprocessing: An Overview
 - Data Quality
 - Major Tasks in Data Preprocessing
- Data Cleaning
- Data Integration 
- Data Reduction
- Data Transformation
- Data Discretization

Data Integration



- Data integration:
 - Combines data from multiple sources into a coherent store
- Schema integration: e.g., $A.\text{cust-id} \equiv B.\text{cust-}\#$
 - Integrate metadata from different sources
- Entity identification problem:
 - Identify real world entities from multiple data sources, e.g.,
Bill Clinton = William Clinton
- Detecting and resolving data value conflicts
 - For the same real world entity, attribute values from different sources are different
 - Possible reasons: different representations, different scales, e.g., metric vs. British units

Handling Redundancy in Data Integration

- Redundant data occur often when integrating multiple databases
 - ***Object identification***: The same attribute or object may have different names in different databases
 - ***Derivable data***: One attribute may be a “derived” attribute in another table, e.g., annual revenue
- Redundant attributes may be possible to detect by ***correlation analysis***
- Careful integration of the data from multiple sources may reduce/avoid redundancies and inconsistencies and improve mining speed and quality

Correlation Analysis (Nominal Data)

- **χ^2 (chi-square) test**

$$\chi^2 = \sum \frac{(\text{Observed} - \text{Expected})^2}{\text{Expected}}$$

- The cells that contribute the most to the χ^2 value are those whose actual count is very different from the expected count
- Correlation does not imply causality
 - # of hospitals and # of car-thefts in a city are correlated
 - Both are causally linked to the third variable: population

Chi-Square Calculation: an Example

- Questionnaire among $N=1500$ participants:

	Play chess	Not play chess	Sum (row)
Like science fiction	250	250	500
Not like science fiction	50	950	1000
Sum (column)	300	1200	1500

- Expected results e_{ij} from the **null hypothesis** stating that “preferred reading” and “game favour” are **uncorrelated**):

	Play chess	Not play chess	Sum (row)
Like science fiction	100	400	500
Not like science fiction	200	800	1000
Sum (column)	300	1200	1500

$$e_{ij} = \text{sum}(\text{col } i) \cdot \text{sum}(\text{row } j) / N$$

Chi-Square Calculation: an Example


- χ^2 (chi-square) calculation:

$$\chi^2 = \frac{(250 - 100)^2}{100} + \frac{(50 - 200)^2}{200} + \frac{(250 - 400)^2}{400} + \frac{(950 - 800)^2}{800} = \underline{421.9}$$

- What does it show?

Chi-square Test


- Small deviations are more expected than large deviations:



- Random variable Y has a normal distribution
→ Y^2 has a chi-square distribution (with 1 degree of freedom)

Chi-square Test

- Some percentage of expected deviations is over a threshold



- E.g.: 10% of all χ^2 values are larger than critical value 2.7
- Values can be looked up in a chi-square distribution table

Chi-square Distribution Table

probability level

α	0.5	0.1	0.05	0.02	0.01	0.001
	0.455	2.706	3.841	5.412	6.635	10.827

- $\alpha=0.1$: 90% of values are below critical value 2.706
 - $\alpha=0.05$: 95% of values are below critical value 3.841
 - ...
 - α : significance level
 - Earlier, we have found a value of 421.9
- Our data are extremely unlikely given the null hypothesis

Chi-Square Calculation: an Example

- χ^2 (chi-square) calculation:

$$\chi^2 = \frac{(250-100)^2}{100} + \frac{(50-200)^2}{200} + \frac{(250-400)^2}{400} + \frac{(950-800)^2}{800} = \underline{421.9}$$

- It shows that “preferred reading” and “game favour” are correlated in the group (since χ^2 larger than 10.827, from χ^2 table – a statistical measure for significance of 2x2 table)
- What if all numbers were 10x smaller ($N=150$ participants)?
 $\rightarrow \chi^2 = \dots = 42.19$
- What if all numbers were 50x smaller ($N=30$ participants)?

$$\chi^2 = \frac{(5-2)^2}{2} + \frac{(1-4)^2}{4} + \frac{(5-8)^2}{8} + \frac{(19-16)^2}{16} = 8.4375$$


Chi-Square Calculation: Degrees of Freedom

		Category 1 Levels				Sum (row)
		L1	L2	...	LJ	
Category 2 Levels	L1					
	...					
	LI					
Sum (col.)						N

- If the two categories have several levels (J levels for category 1 and I levels for category 2), then there are more degrees of freedom in which the entries can differ
- $I \times J$ ***contingency table***
- Number of degrees of freedom: $(I-1) \times (J-1)$

Chi-Square Calculation: Degrees of Freedom

- More degrees of freedom make larger deviations probable:



- Sum of k independ. random variables Y_i with normal distrib.
 $\rightarrow \sum_i Y_i^2$ has a χ^2 distribution with k degrees of freedom

Chi-square Distribution Table

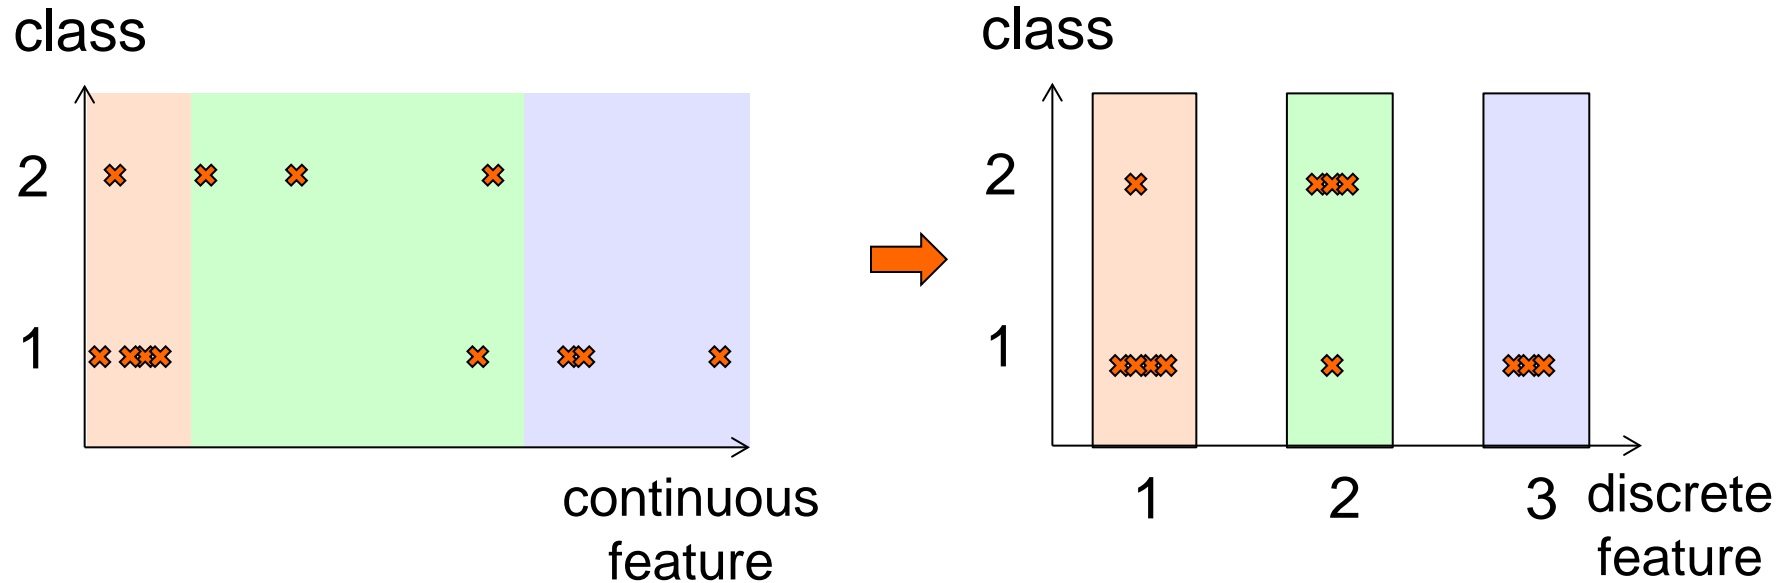
probability level

DF	0.5	0.1	0.05	0.02	0.01	0.001
1	0.455	2.706	3.841	5.412	6.635	10.827
2	1.386	4.605	5.991	7.824	9.210	13.815
3	2.366	6.251	7.815	9.837	11.345	16.268
4	3.357	7.779	9.488	11.668	13.277	18.465
5	4.351	9.236	11.070	13.388	15.086	20.517

- 
- Degrees of freedom

Values Reduction

Data often have continuous features



Many classification algorithms require data with discrete attributes

Values Reduction – ChiMerge Technique

1. **Sort** the data for the given feature in ascending order
 2. **Define initial intervals** so that every value of the feature is in a separate interval
 3. **Repeat**:
 - 3.1 Compute X^2 tests for each pair of adjacent intervals
 - 3.2 Merge two adjacent intervals with the lowest X^2 value, if calculated X^2 is less than threshold
- Until** no X^2 test of any two adjacent intervals is less than threshold value

Values Reduction – Contingency Table

- A ChiMerge requires computation of χ^2 test for the contingency table 2 x 2 of categorical data:

	Class 1	Class 2	Σ
Interval-1	A_{11}	A_{12}	R_1
Interval-2	A_{21}	A_{22}	R_2
Σ	C_1	C_2	N

χ^2 test is:

$$\chi^2 = \sum_{i=1}^2 \sum_{j=1}^k \frac{(A_{ij} - E_{ij})^2}{E_{ij}}$$

where:

k = number of classes,

A_{ij} = number of instances in the i -th interval, j -th class,

E_{ij} = **expected frequency** of A_{ij} , which is computed as $(R_i \cdot C_j) / N$,

R_i = number of instances in the i -th interval = $\sum A_{ij}$, $j = 1, \dots, k$,

C_j = number of instances in the j -th class = $\sum A_{ij}$, $i = 1, 2$,

N = total number of instances = $\sum R_i$, $i = 1, 2$.

Values Reduction – ChiMerge Technique Example

Data Set

Sample: F		K	
1	1	1	0
2	3	2	2
3	7	1	5
4	8	1	7.5
5	9	1	8.5
6	11	2	.
7	23	2	.
8	37	1	.
9	39	2	
10	45	1	
11	46	1	
12	59	1	

Initial interval points

Values Reduction – ChiMerge Technique Example

- χ^2 was minimum for intervals: $[7.5, 8.5]$ and $[8.5, 10]$

Sample: F		K
1	1	1
2	3	2
3	7	1
4	8	1
5	9	1
6	11	2
7	23	2
8	37	1
9	39	2
10	45	1
11	46	1
12	59	1

	Class 1	Class 2	Σ
Interval $[7.5, 8.5]$	$A_{11}=1$	$A_{12}=0$	$R_1=1$
Interval $[8.5, 10]$	$A_{21}=1$	$A_{22}=0$	$R_2=1$
Σ	$C_1=2$	$C_2=0$	$N=2$

Based on the table's values, we can calculate expected values:

$$E_{11} = 1*2/2 = 1, \quad E_{12} = 1*0/2 = 0,$$

$$E_{21} = 1*2/2 = 1, \quad \& \quad E_{22} = 1*0/2 = 0$$

and corresponding χ^2 test:

$$\begin{aligned} \chi^2 &= (1-1)^2/1 + (0-0)^2/0 \\ &\quad + (1-1)^2/1 + (0-0)^2/0 = 0 \end{aligned}$$

For $d=1$ degree of freedom: $\chi^2 = 0 < 2.706 \rightarrow \text{merge !}$
($\alpha=0.1$)

Values Reduction – ChiMerge Technique Example

- ... one of the following iterations:

Sample: F	K	
1	1	1
2	3	2
3	7	1
4	8	1
5	9	1
6	11	2
7	23	2
8	37	1
9	39	2
10	45	1
11	46	1
12	59	1

	Class 1	Class 2	Σ
Interval [0.0 , 7.5]	$A_{11}=2$	$A_{12}=1$	$R_1=3$
Interval [7.5 , 10]	$A_{21}=2$	$A_{22}=0$	$R_2=2$
Σ	$C_1=4$	$C_2=1$	$N=5$

$$E_{11} = 3*4/5 = 2.4, \quad E_{12} = 3*1/5 = 0.6,$$

$$E_{21} = 2*4/5 = 1.6, \quad E_{22} = 2*1/5 = 0.4$$

$$\begin{aligned} X^2 &= (2-2.4)^2/2.4 + (1-0.6)^2/0.6 \\ &\quad + (2-1.6)^2/1.6 + (0-0.4)^2/0.4 = \mathbf{0.834} \end{aligned}$$

For $d=1$ degree of freedom: $X^2 = 0.834 < 2.706 \rightarrow \text{merge !}$
($\alpha=0.1$)

Values Reduction – ChiMerge Technique Example

- ... One of the additional iterations:

Sample: F

K

1	1	1
2	3	2
3	7	1
4	8	1
5	9	1
6	11	2
7	23	2
8	37	1
9	39	2
10	45	1
11	46	1
12	59	1

	Class 1	Class 2	Σ
Interval [0.0 , 10]	$A_{11}=4$	$A_{12}=1$	$R_1=5$
Interval [10 , 42]	$A_{21}=1$	$A_{22}=3$	$R_2=4$
Σ	$C_1=5$	$C_2=4$	$N=9$

E11 = 2.78 , E12 = 2.22 ,

E21 = 2.22 , E22 = 1.78

$X^2 =$

2.72 > 2.706

→ NO merge

$$E_{11} = 2.78, \quad E_{12} = 2.22, \\ E_{21} = 2.22, \quad E_{22} = 1.78$$

$$\chi^2 = 2.72 > 2.706 \quad \rightarrow \text{NO merge !}$$

Final discretization:

[0, 10], [10, 42], and [42, 60]



Interval representatives:

5 (low)



26 (medium)

51 (high)

Values Reduction – ChiMerge Technique Example

Sample: F		K
1	1	1
2	3	2
3	7	1
4	8	1
5	9	1
6	11	2
7	23	2
8	37	1
9	39	2
10	45	1
11	46	1
12	59	1


Final data set
with reduced
set of values F:



Original set


Sample: F		K
1	5	1
2	5	2
3	5	1
4	5	1
5	5	1
6	26	2
7	26	2
8	26	1
9	26	2
10	51	1
11	51	1
12	51	1

Data Preprocessing

- Data Preprocessing: An Overview
 - Data Quality
 - Major Tasks in Data Preprocessing
- Data Cleaning
- Data Integration
- Data Reduction 
- Data Transformation
- Data Discretization

Feature Reduction

- Which features to select, and how?




TRS_DT	TRS_TYP_CD	REF_DT	REF_NUM	CO_CD	GDS_CD	QTY	UT_CD	UT_PRIC
21/05/93	00001	04/05/93	25119	10002J	001M	10	CTN	22.000
21/05/93	00001	05/05/93	25124	10002J	032J	200	DOZ	1.370
21/05/93	00001	05/05/93	25124	10002J	033Q	500	DOZ	1.000
21/05/93	00001	13/05/93	25217	10002J	024K	5	CTN	21.000
21/05/93	00001	13/05/93	25216	10026H	006C	20	CTN	69.000
21/05/93	00001	13/05/93	25216	10026H	008Q	10	CTN	114.000
21/05/93	00001	14/05/93	25232	10026H	006C	10	CTN	69.000
21/05/93	00001	14/05/93	25235	10027E	003A	5	CTN	24.000
21/05/93	00001	14/05/93	25235	10027E	001M	5	CTN	24.000
21/05/93	00001	22/04/93	24974	10035E	009F	50	CTN	118.000
21/05/93	00001	27/04/93	25033	10035E	015A	375	GRS	72.000
21/05/93	00001	20/05/93	25313	10041Q	010F	10	CTN	26.000
21/05/93	00001	12/05/93	25197	10054R	002E	25	CTN	24.000

Features Reduction

Two standard approaches:

- ***Feature selection***: A process that chooses an optimal subset of features according to an objective function:
 - feature ranking algorithms,
 - minimum subset algorithms.
- ***Feature extraction***: refers to the mapping of the original high-dimensional data onto a lower-dimensional space.
 - Descriptive setting: minimizes the information loss
 - Predictive setting: maximizes the class discrimination

Feature selection – Example for Optimal Features' Subset




F_1	F_2	F_3	F_4	F_5	C
0	0	1	0	1	0
0	1	0	0	1	1
1	0	1	0	1	1
1	1	0	0	1	1
0	0	1	1	0	0
0	1	0	1	0	1
1	0	1	1	0	1
1	1	0	1	0	1

- Data set (whole set)
 - Five Boolean features
 - $C = F_1 \vee F_2$
 - $F_3 = \neg F_2$, $F_5 = \neg F_4$
 - Optimal subset:
 $\{F_1, F_2\}$ or $\{F_1, F_3\}$
- Combinatorial nature of
searching for an optimal subset

Feature Selection – Complexity

- **Feature selection** in general can be viewed as a search problem.
- For practical methods, an optimal search is not feasible, and simplifications are made to produce acceptable and timely reasonable results:

- heuristic criteria
- bottom-up approach
- top-down approach



Methods of Feature Selection

- Univariate methods
 - Considers one variable (feature) at a time.
- Filter methods
 - Separating feature selection from classifier learning
 - Relying on general characteristics of data (information, distance, dependence, consistency)
 - No bias toward any learning algorithm, fast
- Wrapper methods
 - Relying on a predetermined classification algorithm.
 - Using predictive accuracy as goodness measure
 - High accuracy, computationally expensive
- Embedded methods
 - Combine Filter and Wrapper approaches

Features Selection: Univariate Methods

Comparison of means and variances:

- Samples of two classes (A and B) can be examined:

$$SE(A - B) = \sqrt{\frac{\text{var}(A)}{n_1} + \frac{\text{var}(B)}{n_2}}$$

- TEST:
$$\frac{|\text{mean}(A) - \text{mean}(B)|}{SE(A - B)} > \text{threshold-value}$$

where n_1 and n_2 are the corresponding number of samples for classes A and B .

Features Selection: Univariate Methods

- Comparison of *means* and *variances* – **Example:**

X	Y	C
0.3	0.7	A
0.2	0.9	B
0.6	0.6	A
0.5	0.5	A
0.7	0.7	B
0.4	0.9	B

Threshold value is 0.5

$$X_A = \{0.3, 0.6, 0.5\},$$

$$X_B = \{0.2, 0.7, 0.4\},$$

$$Y_A = \{0.7, 0.6, 0.5\},$$

$$Y_B = \{0.9, 0.7, 0.9\}$$

Features Selection: Univariate Methods

- Comparison of *means* and *variances* – **Example:**

$$SE(X_A - X_B) = \sqrt{\frac{\text{var}(X_A)}{n_1} + \frac{\text{var}(X_B)}{n_2}} = \sqrt{\frac{0.0233}{3} + \frac{0.6333}{3}} = 0.4678$$

$$SE(Y_A - Y_B) = \sqrt{\frac{\text{var}(Y_A)}{n_1} + \frac{\text{var}(Y_B)}{n_2}} = \sqrt{\frac{0.01}{3} + \frac{0.0133}{3}} = 0.0875$$


Tests:

$$X: \frac{|\text{mean}(A) - \text{mean}(B)|}{SE(A - B)} = \frac{|0.4667 - 0.4333|}{0.4678} < 0.5$$

$$Y: \frac{|\text{mean}(A) - \text{mean}(B)|}{SE(A - B)} = \frac{|0.6 - 0.8333|}{0.0875} > 0.5$$

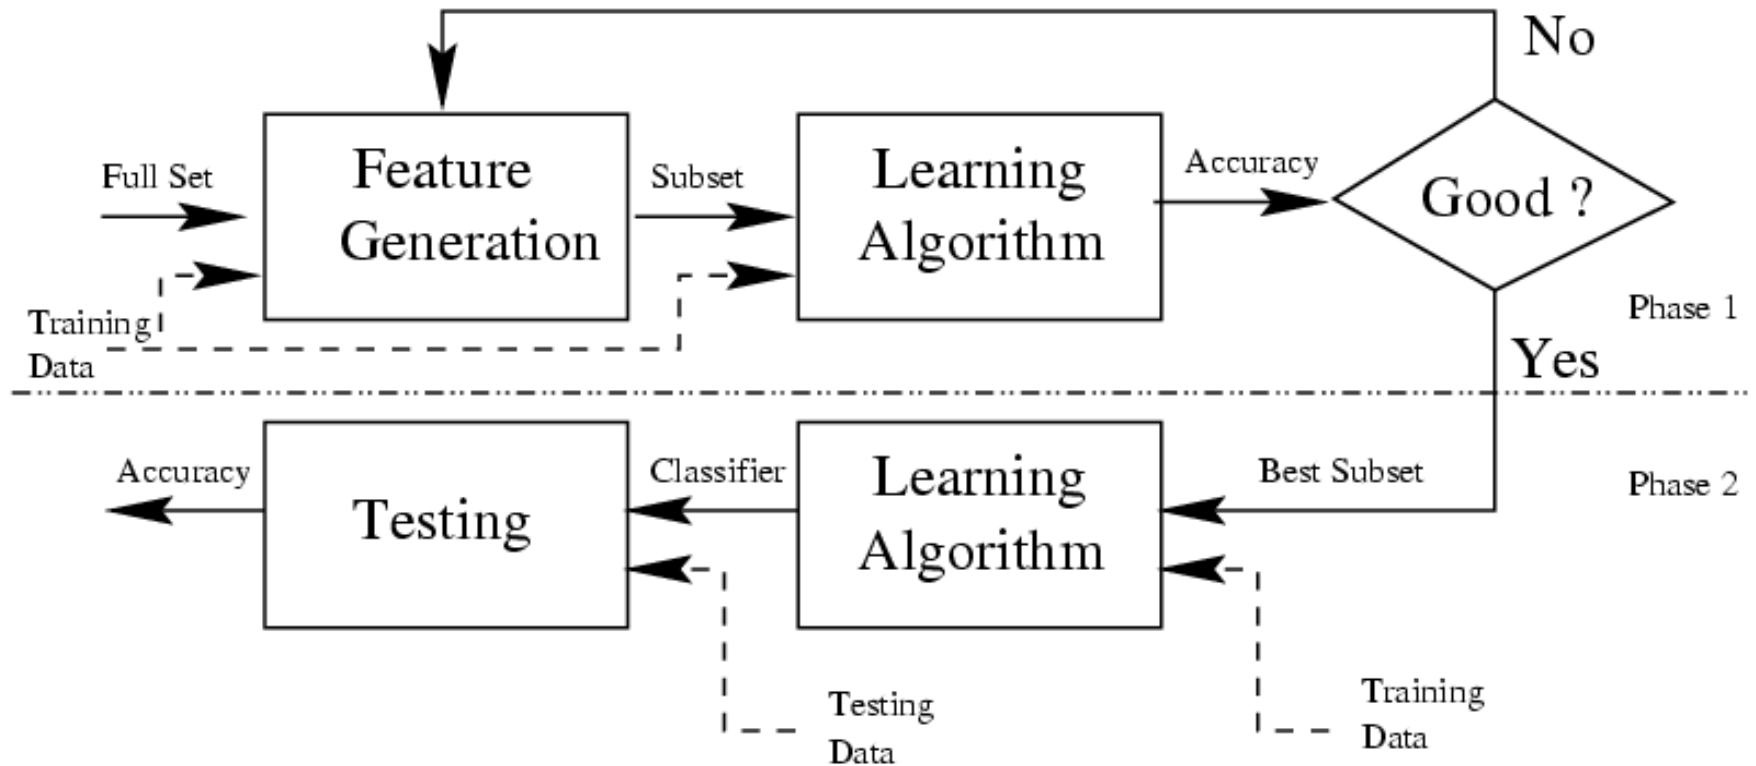
X is a candidate feature for reduction because its mean values are close, and therefore the final test is below threshold value.

Filter Model



- Example filter algorithm for feature selection:
 - **Relief** (Kira & Rendell 1992)

Wrapper Model



- Example wrapper algorithm for Feature Selection:
 - **SVM**

Data Reduction: Dimensionality Reduction

■ ***Curse of dimensionality***

- When dimensionality increases, data becomes increasingly sparse
- Density and distance between points, which is critical to clustering, outlier analysis, becomes less meaningful
- The possible combinations of subspaces will grow exponentially

■ ***Dimensionality reduction***

- Avoid the curse of dimensionality
- Help eliminate irrelevant features and reduce noise
- Reduce time and space required in data mining
- Allow easier visualization


Data Reduction Strategies

- Why data reduction? — A database/data warehouse may store terabytes of data. Complex data analysis may take a very long time to run on the complete data set.
- Obtain a reduced representation of the data set that is much smaller in volume but yet produces the (almost) same analytical results
- Data reduction strategies
 - **Dimensionality reduction**, e.g., remove unimportant attributes
 - Feature subset selection, feature creation
 - Wavelet transforms
 - Principal Components Analysis (PCA)
 - **Numerosity** reduction (some simply call it: Data Reduction)
 - Regression and Log-Linear Models
 - Histograms, clustering, sampling
 - Data cube aggregation
 - **Data compression**


Dimensions Reduction of Large Data Sets

Main dimensions:

- **columns** (features),
- **rows** (cases or samples),
- **values** of the features for the given sample




Data Preprocessing

- Data Preprocessing: An Overview
 - Data Quality
 - Major Tasks in Data Preprocessing
- Data Cleaning
- Data Integration
- Data Reduction
- Data Transformation 
- Data Discretization


Mapping Data to a New Space

- Fourier transform: mapping from time to frequency domain
- Wavelet transform



Two Sine Waves

Two Sine Waves + Noise



Frequencies

Data Transformation

- A function that maps the entire set of values of a given attribute to a new set of replacement values so that each old value can be identified with one of the new values
- Methods
 - Smoothing: Remove noise from data
 - Attribute/feature construction
 - New attributes constructed from the given ones
 - Aggregation: Summarization
 - Normalization: Scaled to fall within a smaller, specified range
 - min-max normalization
 - z-score normalization
 - normalization by decimal scaling
 - Discretization: Concept hierarchy climbing

Normalization

- **Min-max normalization**: to $[\text{new_min}_A, \text{new_max}_A]$

$$v' = \frac{v - \text{min}_A}{\text{max}_A - \text{min}_A} (\text{new_max}_A - \text{new_min}_A) + \text{new_min}_A$$

- **Ex.** Let income range \$12,000 to \$98,000. Normalize to [0.0, 1.0].

Then \$73,600 is mapped to $\frac{73,600 - 12,000}{98,000 - 12,000} (1.0 - 0) + 0 = 0.716$

- **Z-score normalization** (μ : mean, σ : standard deviation):

$$v' = \frac{v - \mu_A}{\sigma_A}$$

- **Ex.** Let $\mu = 54,000$, $\sigma = 16,000$. Then $\frac{73,600 - 54,000}{16,000} = 1.225$

- **Normalization by decimal scaling**

$$v' = \frac{v}{10^j} \quad \text{where } j \text{ is the smallest integer such that } \text{Max}(|v'|) < 1$$

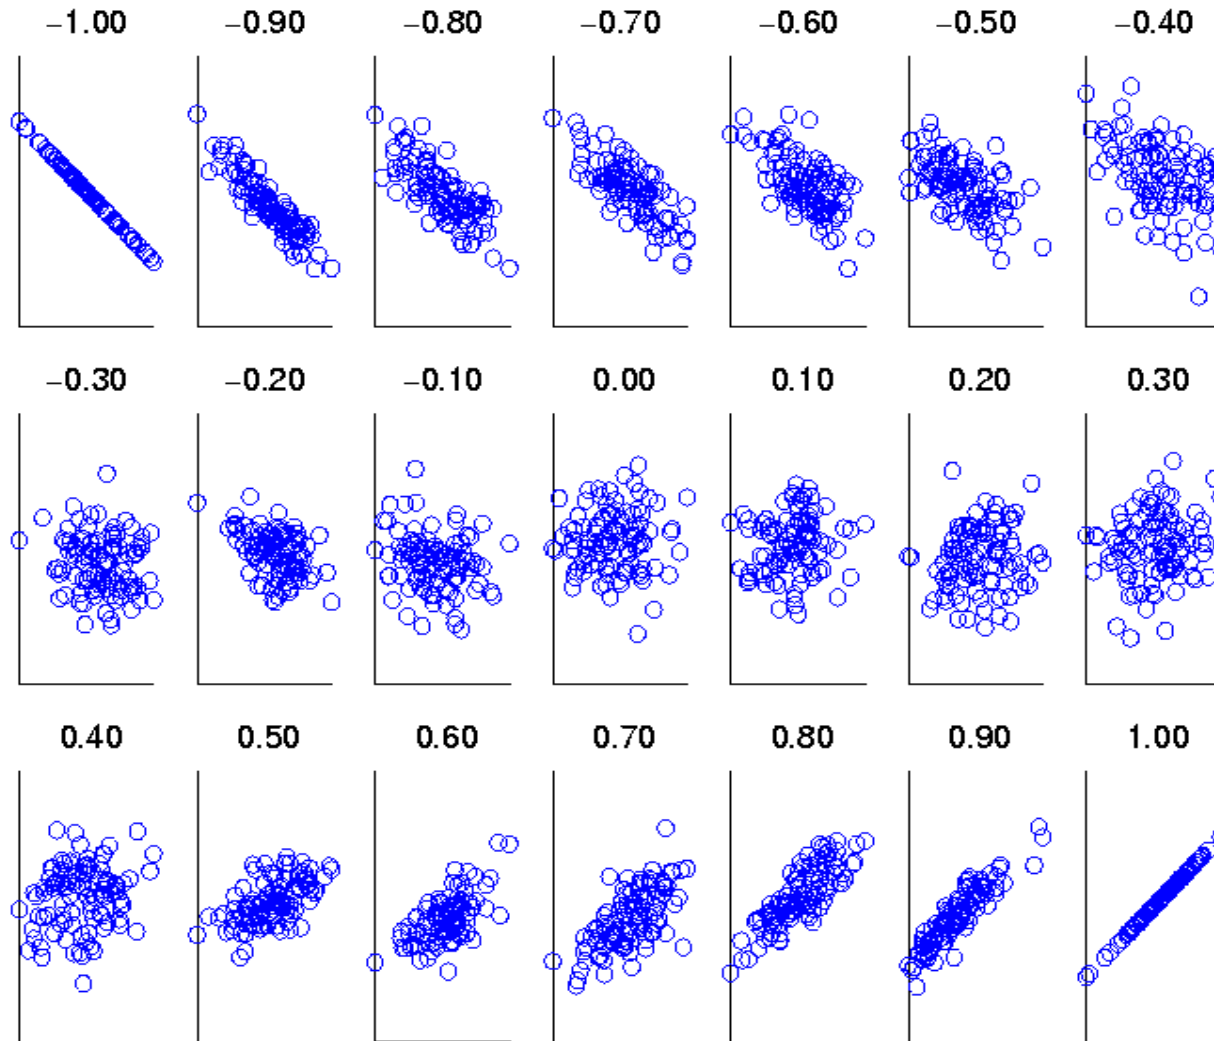
Correlation Analysis (Numeric Data)

- Correlation coefficient (also called ***Pearson's product moment coefficient***)

$$r_{A,B} = \frac{\sum_{i=1}^N (a_i - \bar{A})(b_i - \bar{B})}{N\sigma_A\sigma_B}$$

- N = number of tuples
 - \bar{A} and \bar{B} = means of attributes A and B
 - σ_A and σ_B = standard deviations of A and B
- $r_{A,B} > 0$: A and B are positively correlated
 - A's values increase as B's. Larger $r_{A,B} \rightarrow$ stronger correlation.
- $r_{A,B} = 0$: independent
- $r_{AB} < 0$: negatively correlated

Visually Evaluating Correlation



Scatter plots
Correlation
coefficients range
from -1 to 1.

Covariance (Numeric Data)

- Covariance:

$$Cov(A, B) = E((A - \bar{A})(B - \bar{B})) = \frac{\sum_{i=1}^n (a_i - \bar{A})(b_i - \bar{B})}{n}$$

Related to correlation coefficient: $r_{A,B} = \frac{Cov(A, B)}{\sigma_A \sigma_B}$

- n = number of tuples
 - \bar{A} and \bar{B} = mean or **expected values** of A and B
 - σ_A and σ_B = standard deviation of A and B .
- **Positive covariance:** If $Cov_{A,B} > 0$, then A and B both tend to be larger than their expected values.
 - **Negative covariance:** If $Cov_{A,B} < 0$ then if A is larger than its expected value, B is likely to be smaller than its expected value.
 - **Independence:** $Cov_{A,B} = 0$

Co-Variance: an Example

$$Cov(A, B) = E((A - \bar{A})(B - \bar{B})) = \frac{\sum_{i=1}^n (a_i - \bar{A})(b_i - \bar{B})}{n}$$

- It can be simplified in computation as

$$Cov(A, B) = E(A \cdot B) - \bar{A}\bar{B}$$

- Suppose two stocks A and B have the following values in one week: (2, 5), (3, 8), (5, 10), (4, 11), (6, 14).
- **Question:** If the stocks are affected by the same industry trends, will their prices rise or fall together?
 - $E(A) = (2 + 3 + 5 + 4 + 6) / 5 = 20/5 = 4$
 - $E(B) = (5 + 8 + 10 + 11 + 14) / 5 = 48/5 = 9.6$
 - $Cov(A, B) = (2 \cdot 5 + 3 \cdot 8 + 5 \cdot 10 + 4 \cdot 11 + 6 \cdot 14) / 5 - 4 \cdot 9.6 = 4$
- Thus, A and B rise together since $Cov(A, B) > 0$.

Principal Component Analysis (Steps)

- Given N data vectors from n -dimensions, find $m \ll n$ orthogonal vectors (***principal components***) to represent the data
 - Subtract mean from the input data, each attribute has mean zero
 - Compute m orthonormal (unit) vectors, i.e., *principal components*
 - Each input data (vector) is a linear combination of the m principal component vectors
 - The principal components are sorted in order of decreasing “***significance***” (1st component: data has maximum variance)
 - Since the components are sorted, the size of the data can be reduced by eliminating the *insignificant components*, i.e., those with low variance
 - Thus, using only the most significant principal components, it is possible to reconstruct a good approximation of the original data.
- Works for numeric data; reduction of higher dimensions to lower

Principal Components Analysis Algorithm

1. Compute the $n \times n$ **covariance matrix** S

$$S_{ij} = \frac{1}{N-1} \cdot \sum_{d=1}^N (x_i^d - \bar{x}_i)^T \cdot (x_j^d - \bar{x}_j) \quad \text{where} \quad \bar{x} = \frac{1}{N-1} \cdot \sum_{d=1}^N x^d$$

2. Calculate the **eigenvalues** of the covariance matrix S for the given data and sort the eigenvalues of S :

$$\{\lambda_1, \lambda_2, \dots, \lambda_n\} \quad \text{where} \quad \lambda_1 \geq \lambda_2 \geq \dots \geq \lambda_n \geq 0.$$


The eigenvalues are the variances of the data in the directions of ``their`` eigenvectors.

3. The **eigenvectors** e_1, e_2, \dots, e_n correspond to ``their`` eigenvalues $\lambda_1, \lambda_2, \dots, \lambda_n$,

The eigenvectors are called the **principal axes**.


Principal Component Analysis (PCA)

- Find a projection that captures the largest amount of variation in data
- The original data are projected onto a much smaller space, resulting in dimensionality reduction. We find the eigenvectors of the covariance matrix, and these eigenvectors define the new space




Multiply a Vector with an nxn Matrix: $A \cdot x$

- Symmetric matrix A




Multiply a Vector with an nxn Matrix: $A \cdot x$

- Non-symmetric matrix A




Multiply a Vector with an nxn Matrix: $A \cdot x$

- Arbitrary matrix A



Multiply a Vector with an $n \times n$ Matrix: $A \cdot x$


- Rotation matrix A




No eigenvectors

Multiply a Vector with a Covariance Matrix


- The covariance matrix S is a symmetric matrix
 - It always has **eigenvectors**
 - Eigenvectors are **orthogonal** to each other, if their eigenvalues differ
- Its **eigenvalues** are ≥ 0
 - S is positive (semi-) definite
- Eigenvalues are **variances** of the data in the directions of the corresponding eigenvectors
- Multiplying by S does not rotate any vector by more than 90°



Repeatedly Multiply with the Covariance Matrix



Repeatedly Multiply with the Covariance Matrix



New coordinate system:

Axes defined by the eigenvectors e_i

Here: $x \rightarrow \tilde{x}$

Here, covariance matrix expressed as a diagonal matrix D

The diagonal elements of D are the eigenvalues λ_i

Easy to see: individual dimensions/components i of \tilde{x} are multiplied by λ_i

Obtain Largest Eigenvalue and Eigenvector

Iterative method

- Choose an initial random vector x
- Repeat
 - $x \leftarrow S \cdot x$
 - Normalize x to length 1
- Until converged. Then **normalized eigenvector** $e_1 = x$

Compute corresponding **eigenvalue** as the norm:

$$\lambda_1 = \|S \cdot e_1\|$$

- Other linear algebra methods, e.g. SVD, compute **all** eigenvectors and corresponding eigenvalues.


Principal Components Analysis

- The data can be expressed in the new coordinate system

$$x = \bar{x} + \sum_{j=1}^m w_j \cdot e_j$$

- w_j are the data coordinates along the eigenvector axes
- $m < n$: data are only approximately reconstructed

The first principal component is an axis in the direction of **maximum variance**.




Principal Components Analysis

- The criterion for features selection is based on the ratio R of the sum of the m largest eigenvalues ($m \leq n$) of S to the trace of S (for example $R > 90\%$):

$$R = \frac{\sum_{i=1}^m \lambda_i}{\sum_{i=1}^n \lambda_i}$$

← sum over all explained variances

Trace of S
= sum over all variances




- Benefit:
 - non-explained variance $\sum_{i=m+1}^n \lambda_i$ may be small, even if $m \ll n$.

Principle Components Example: Eigenfaces

Eigenfaces are the (first few) principle components of a face image database.

reconstructed
face →

eigenfaces →



Sampling


- Sampling: obtaining a small sample s to represent the whole data set N
- Allow a mining algorithm to run in complexity that is potentially sub-linear to the size of the data
- Key principle: Choose a **representative** subset of the data
 - Simple random sampling may have very poor performance in the presence of skewed data
 - Develop adaptive sampling methods, e.g., stratified sampling

Cases Reduction: Sampling ...


Key principle for effective sampling:

- Using a sample will work almost as well as using the entire data sets, if the sample is *representative*.
- A sample is representative if it has approximately the same property (of interest) as the original set of data.

Cases Reduction: Sample Size



8000 points



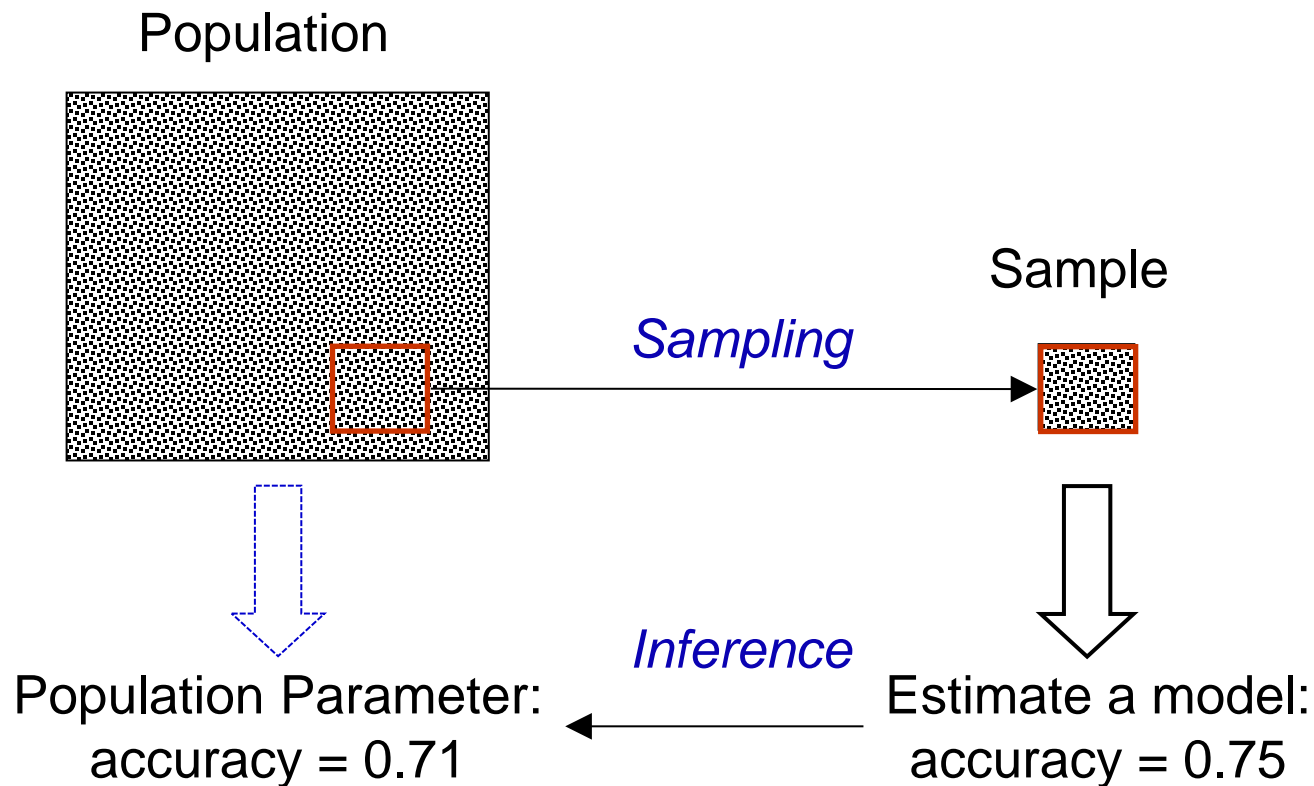
2000 Points



500 Points

Cases Reduction: Accuracy Parameter Estimation

- **Challenging task:** Infer the value of a population parameter based on a sample model.




General-purpose Sampling Methods

■ ***Systematic sampling:***


- For example 50% of a data set (every second sample)
- Simplest.
- Built in most of Data Mining tools
- Problem: regularities in data set!

■ ***Random sampling***

- Random sampling without replacement,
- Random sampling with replacement.
- Average sampling: Combined solution from several subsets (randomly selected)
 - Stratified sampling



Sampling With or Without Replacement




Once an object is selected, it is removed from the population


A selected object is not removed from the population

Stratified Sampling

- Partition the data set into strata (non-overlapping)
- Draw samples from each partition proportionally to its percentage in the data



Data Preprocessing

- Data Preprocessing: An Overview
 - Data Quality
 - Major Tasks in Data Preprocessing
- Data Cleaning
- Data Integration
- Data Reduction
- Data Transformation
- Data Discretization 

Data Discretization Methods

- Reduce number of values for given continuous attribute by dividing into intervals
 - **Binning**: equal width binning and replacing bin by mean
 - Top-down split, unsupervised, no class information used
 - **Histogram analysis**
 - Top-down split, unsupervised, no class information used
 - **Clustering analysis** (unsupervised, top-down split or bottom-up merge)
 - **Decision-tree analysis** (supervised, top-down split)
 - **Correlation** (e.g., χ^2) **analysis** (unsupervised, bottom-up merge)

Simple Discretization: Binning

- **Equal-width** (distance) partitioning
 - Divides the range into N intervals of equal size: uniform grid
 - If A and B are the lowest and highest values of the attribute, the width of intervals will be: $W = (B - A)/N$.
 - Simple, but outliers may dominate presentation
 - Skewed data is not handled well
- **Equal-depth** (frequency) partitioning
 - Divides the range into N intervals, each containing approximately same number of samples
 - Good data scaling

Binning Methods for Data Smoothing

Sorted data for price (in dollars): 4, 8, 9, 15, 21, 21, 24, 25, 26, 28, 29, 34

Partition into equal-frequency (*equi-depth*) bins:

- Bin 1: 4, 8, 9, 15
- Bin 2: 21, 21, 24, 25
- Bin 3: 26, 28, 29, 34


Smoothing by ***bin means***:

- Bin 1: 9, 9, 9, 9
- Bin 2: 23, 23, 23, 23
- Bin 3: 29, 29, 29, 29

Smoothing by ***bin boundaries***:


- Bin 1: 4, 4, 4, 15
- Bin 2: 21, 21, 21, 25
- Bin 3: 26, 26, 26, 34

Example: Data Resampling and Smoothing in Point Cloud Application




Clustering

- Partition data set into clusters based on similarity, and store cluster representation (e.g., centroid and diameter) only
- Can have hierarchical clustering and be stored in multi-dimensional index tree structures
- There are many choices of clustering definitions and clustering algorithms



Automatic Concept Hierarchy Generation

- Some hierarchies can be automatically generated based on the analysis of the number of distinct values per attribute in the data set
 - The attribute with the most distinct values is placed at the lowest level of the hierarchy
 - Exceptions, e.g., weekday, month, quarter, year



Time-dependent Data

- The **time series** of values can be expressed as a list:
$$X = \{t(1), t(2), t(3), \dots, t(n)\},$$

where $t(n)$ is the most recent value.
- For many problems based on time series the goal is to:
 - **forecast** $t(n+1)$ from previous n values of the feature (or more general forecast $t(n+j)$), where these values are directly related to the predicted value, or
 - **find patterns** in time series.
- The most important step in preprocessing of row time-dependent data is specification of a window or a time lag

Time-dependent Data

- **Example:** time series consisting of eleven measurements:

$\mathbf{X} = \{t(0), t(1), t(2), t(3), t(4), t(5), t(6), t(7), t(8), t(9), t(10)\}$

- window size:

$w=5$

- next value:

$j=1$

Sample	W I N D O W					Next Value
	M1	M2	M3	M4	M5	
1	<u>t(0)</u>	t(1)	t(2)	t(3)	t(4)	t(5)
2	<u>t(1)</u>	t(2)	t(3)	t(4)	t(5)	t(6)
3	<u>t(2)</u>	t(3)	t(4)	t(5)	t(6)	t(7)
4	<u>t(3)</u>	t(4)	t(5)	t(6)	t(7)	t(8)
5	<u>t(4)</u>	t(5)	t(6)	t(7)	t(8)	t(9)
6	<u>t(5)</u>	t(6)	t(7)	t(8)	t(9)	t(10)

Time-dependent Data

- **Example:** time series consisting of eleven measurements:

$\mathbf{X} = \{t(0), t(1), t(2), t(3), t(4), t(5), t(6), t(7), t(8), t(9), t(10)\}$

- window size:

$w=5$

- next value:

$j=3$

Sample	W	I	N	D	O	W	Next Value
	M1	M2	M3	M4	M5		
1	<u>t(0)</u>	t(1)	t(2)	t(3)	t(4)		t(7)
2	<u>t(1)</u>	t(2)	t(3)	t(4)	t(5)		t(8)
3	<u>t(2)</u>	t(3)	t(4)	t(5)	t(6)		t(9)
4	<u>t(3)</u>	t(4)	t(5)	t(6)	t(7)		t(10)

Time-dependent Data

Time-dependent **2D** data

Time	a	b
1	5	117
2	8	113
3	4	116
4	9	118
5	10	119
6	12	120




Samples prepared for **window** $w = 3$

Sample	a (n-2)	a (n-1)	a(n)	b (n-2)	b (n-1)	b(n)
1	5	8	4	117	113	116
2	8	4	9	113	116	118
3	4	9	10	116	118	119
4	9	10	12	118	119	120

Time-dependent Data

- One way of **summarizing** features in the data set is to average them, producing “**moving averages**” (MA):

$$MA(i, m) = \frac{1}{m} \cdot \sum_{j=i-m+1}^i t(j)$$




- The objective is to **smooth** neighboring time points by a moving average to reduce the random variation and noise components:

$$MA(i, m) = t(i) = \text{mean}(i) + \text{error}$$

Spatial-Temporal Data

Jan



Average monthly
temperature

New disciplines: Temporal, Spatial, and Streaming Data Mining

Noise example in real world from the WTM lab

www.informatik.uni-hamburg.de/WTM or
www.knowledge-technology.info



Summary

- **Data quality**: accuracy, completeness, consistency, timeliness, believability, interpretability
- **Data cleaning**: e.g. missing/noisy values, outliers
- **Data integration** from multiple sources:
 - Entity identification problem; Remove redundancies; Detect inconsistencies
 - Chi-square for correlation analysis
- **Data reduction**
 - Dimensionality reduction, e.g. by comparison of means and variances; Numerosity reduction; Data compression
- **Data transformation**
 - Normalization
 - PCA
- **Data discretization**
 - Binning, Clustering, Concept hierarchy generation

HiWi Position at WTM Available

- 20 - 40 hours (one-off, approx. €400-€500)
- Immediate availability
- Task: Organize the Publications on the WTM Web Page