


Data Mining

Lecture 4 Decision Trees and Classification



<http://www.informatik.uni-hamburg.de/WTM/>

Overview

- Decision trees for classification
- Decision tree induction
- Criteria for attribute split
 - Information Gain
 - Gini Impurity
- Continuous attributes
- Gain ratio
- Missing values
- Pruning
- Limitations of decision trees


Motivation: Making decisions

One of 900+
TEDTalks


New ideas every weekday
TED.com

Dan Gilbert: Why we make bad decisions,
TED talks,. [Video online](#)

Decision Boundaries



Decision Boundaries



History of Decision Trees

- 1966: Hunt, colleagues in psychology used full search decision tree methods to model human concept learning
- 1977: Breiman, Friedman, colleagues in statistics develop simultaneous Classification And Regression Trees (CART)
- 1986: Quinlan's landmark paper on ID3
- Late 1980s: Various improvements, i.e: coping with noise, continuous attributes, missing data, non-axis-parallel DTs
- 1993: Quinlan's updated algorithm, C4.5
- Towards 2000: Quinlan: More pruning, overfitting control heuristics (C5.0, etc.); combining DTs; incremental learning

Supervised vs. Unsupervised Learning (discrete outputs)

today

■ Supervised Learning (*Classification*)

- The training data (observations, measurements, etc.) are accompanied by ***categorical class labels*** (discrete or nominal)
- New data is classified based on the training

■ Unsupervised Learning (*Clustering*)

- Given a set of measurements, observations, etc. of which the class labels are unknown
- Aim: establishing the existence of clusters in the data

Supervised vs. Unsupervised Learning (continuous outputs)

■ Supervised Learning (*Regression*)

- The training data (observations, measurements, etc.) are accompanied by ***continuous output values***
- For new data where output values are missing, “predict” the most likely output values based on the training

■ Unsupervised Learning (*general case*)

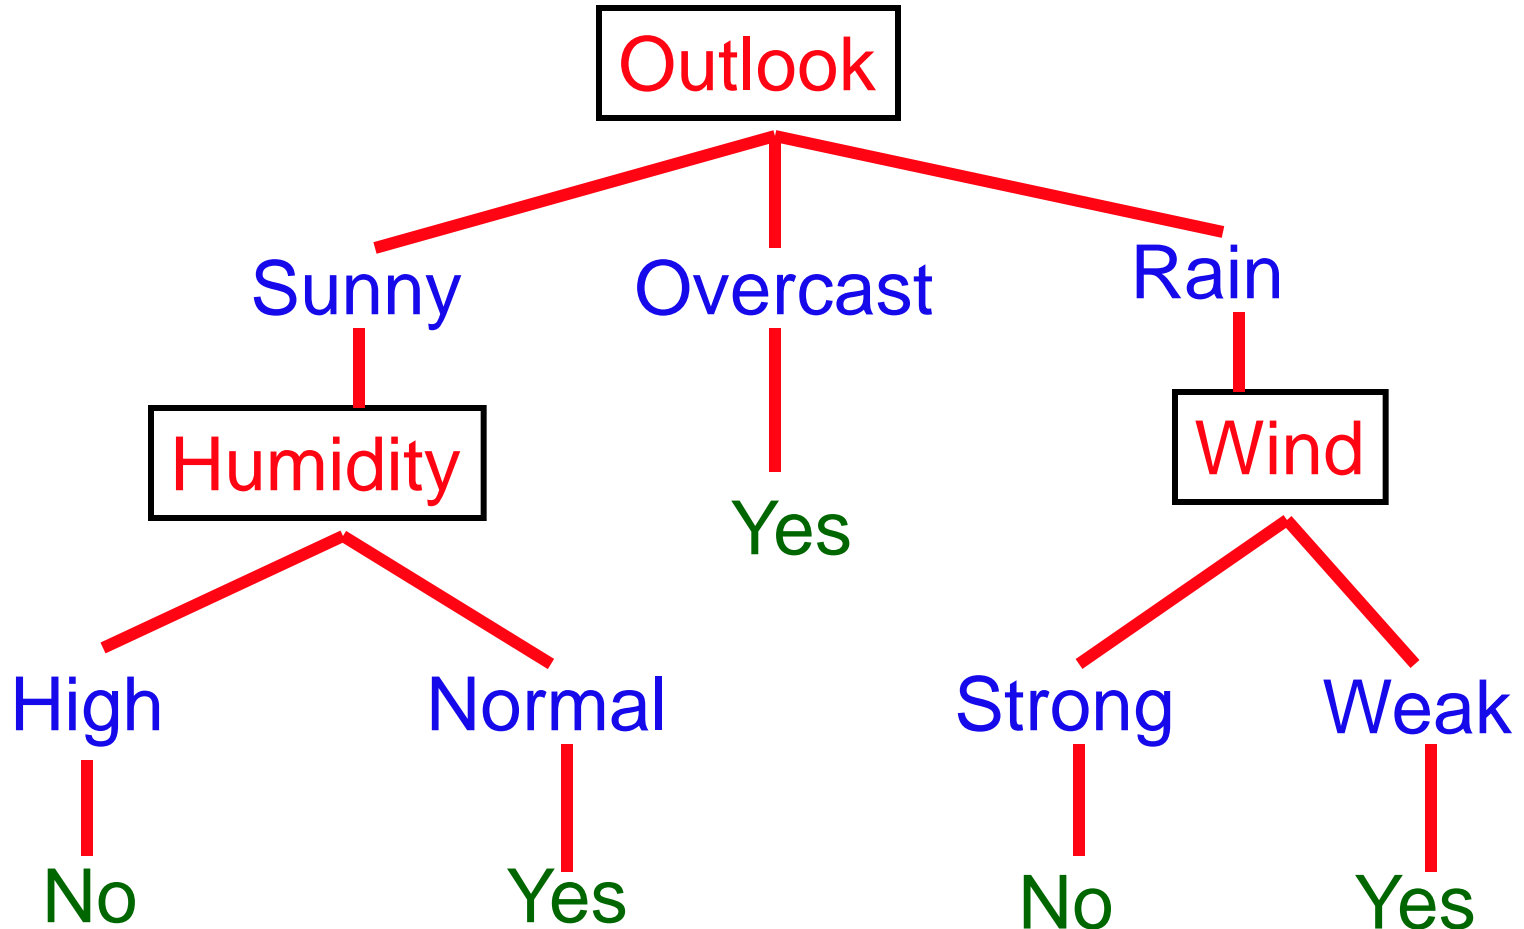
- Given a set of measurements, observations, etc. of which the class labels are unknown
- Aim: represent the data in another form
 - E.g. compressed via PCA, ...

Decision Trees

- Split *classification* into a series of choices about features in turn
- Lay them out in a tree
- Progress down the tree to the leaves

Example: Anyone for Tennis?

Decide whether to play tennis based on the weather




Rules and Decision Trees

- Tree can be turned into a set of rules:
 - if (outlook = sunny & humidity = normal) | (outlook = overcast)
| (outlook = rain & wind = weak)
then
play tennis
- How do we generate the trees?
 - Need to choose features / attributes
 - Need to choose order of features / attributes

Decision Trees


- Efficient method for producing **classifiers** from data
 - **Supervised learning** methods that construct decision trees from a set of input-output samples.
 - Guarantees that a simple tree is found
 - but not necessarily the simplest one
- Consists of
 - **Nodes** that are tests on the attributes
 - Outgoing **branches** of a node correspond to all the possible outcomes of the test at the node
 - **Leaves** that are sets of samples belonging to the same class



Example of Decision Tree for Credit Approval

Credit Analysis

salary	education	label	#
10000	high school	reject	1
40000	under graduate	accept	2
18000	under graduate	reject	3
75000	graduate	accept	4
15000	graduate	accept	5




Decision Tree for Classification

- Given:
 - Database of *samples*, each assigned a *class label*.
- Task: Develop a *model*/profile for each class:
 - **Example profile** (good credit):
(salary > 20k) or (salary <= 20k and education = graduate)
=> Credit = Good (approved)


Classification by Decision Tree Induction

- Decision tree **generation** consists of two phases:
 1. Tree **construction**:
 - At start, all the training examples are at the root.
 - Partition the examples recursively based on selected attributes.
 2. Tree **pruning**:
 - Identify and remove branches that reflect noise or outliers.
- Decision tree **use**: Classifying an unknown sample
 - Test the attribute values of the sample against the decision tree

Decision Tree: Example



Note: tree boundaries are piecewise linear and axis-parallel

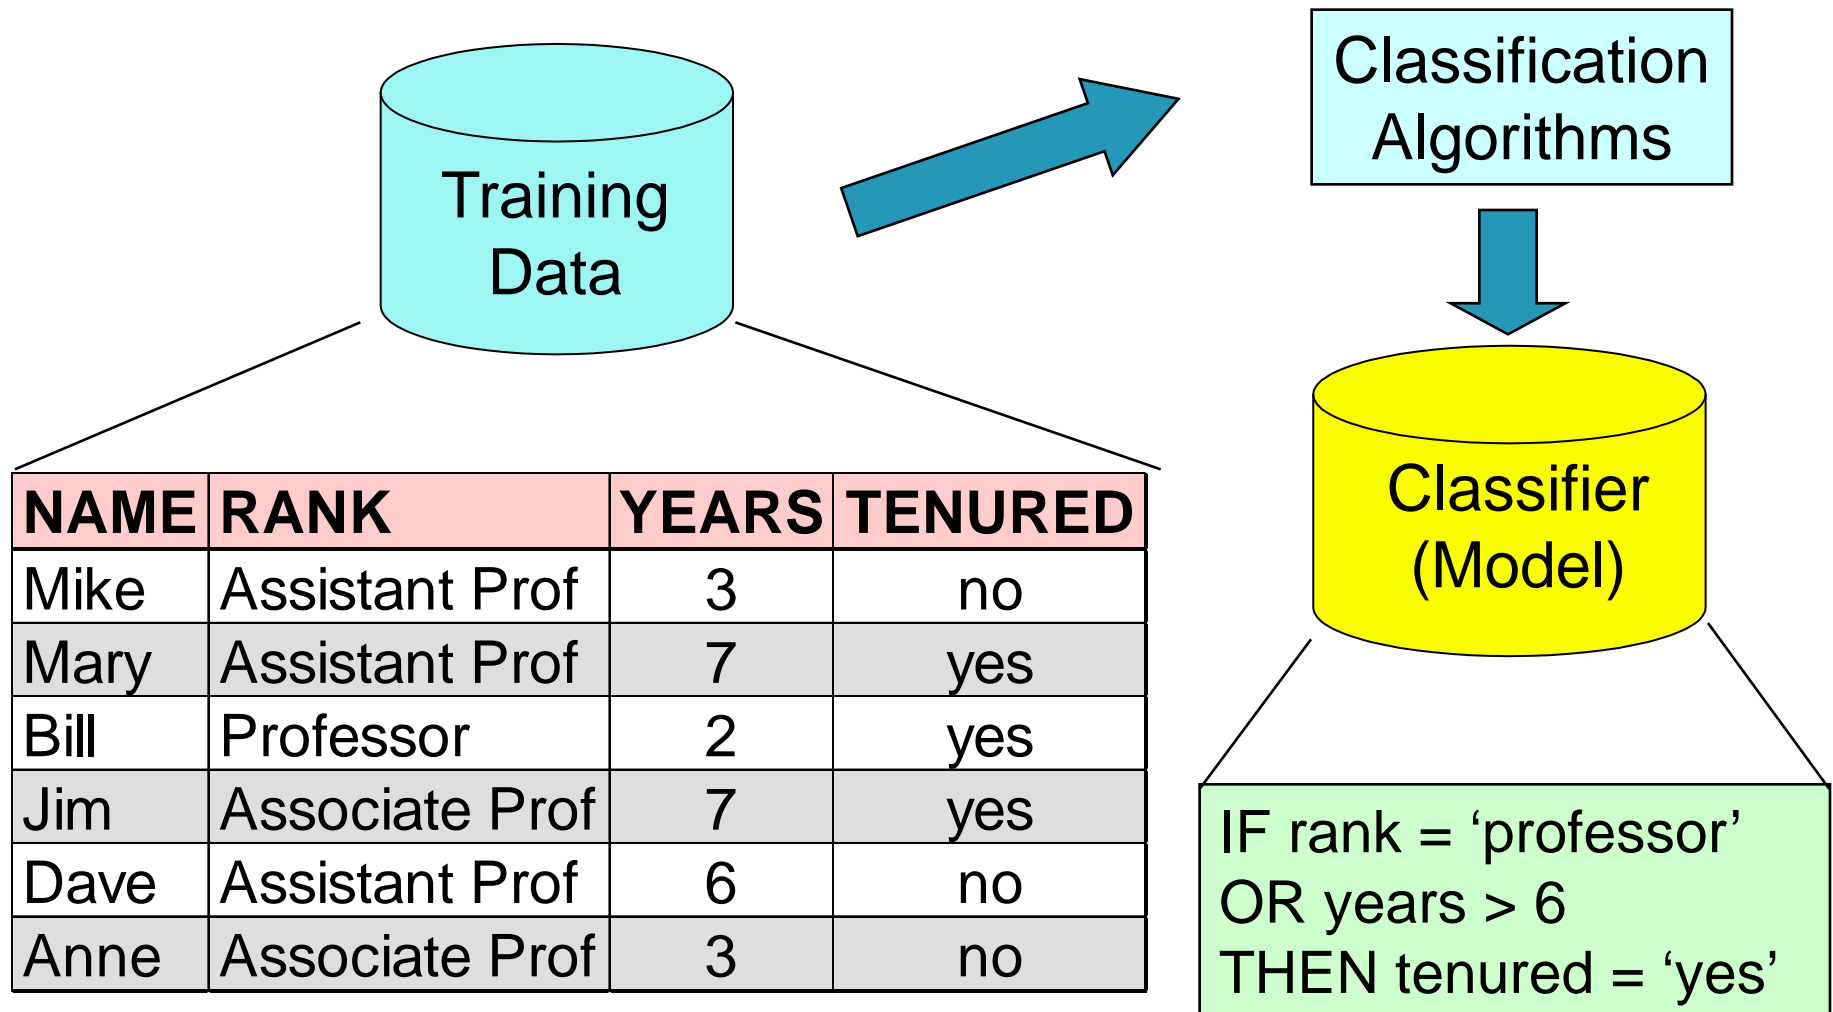


Are all correctly classified?

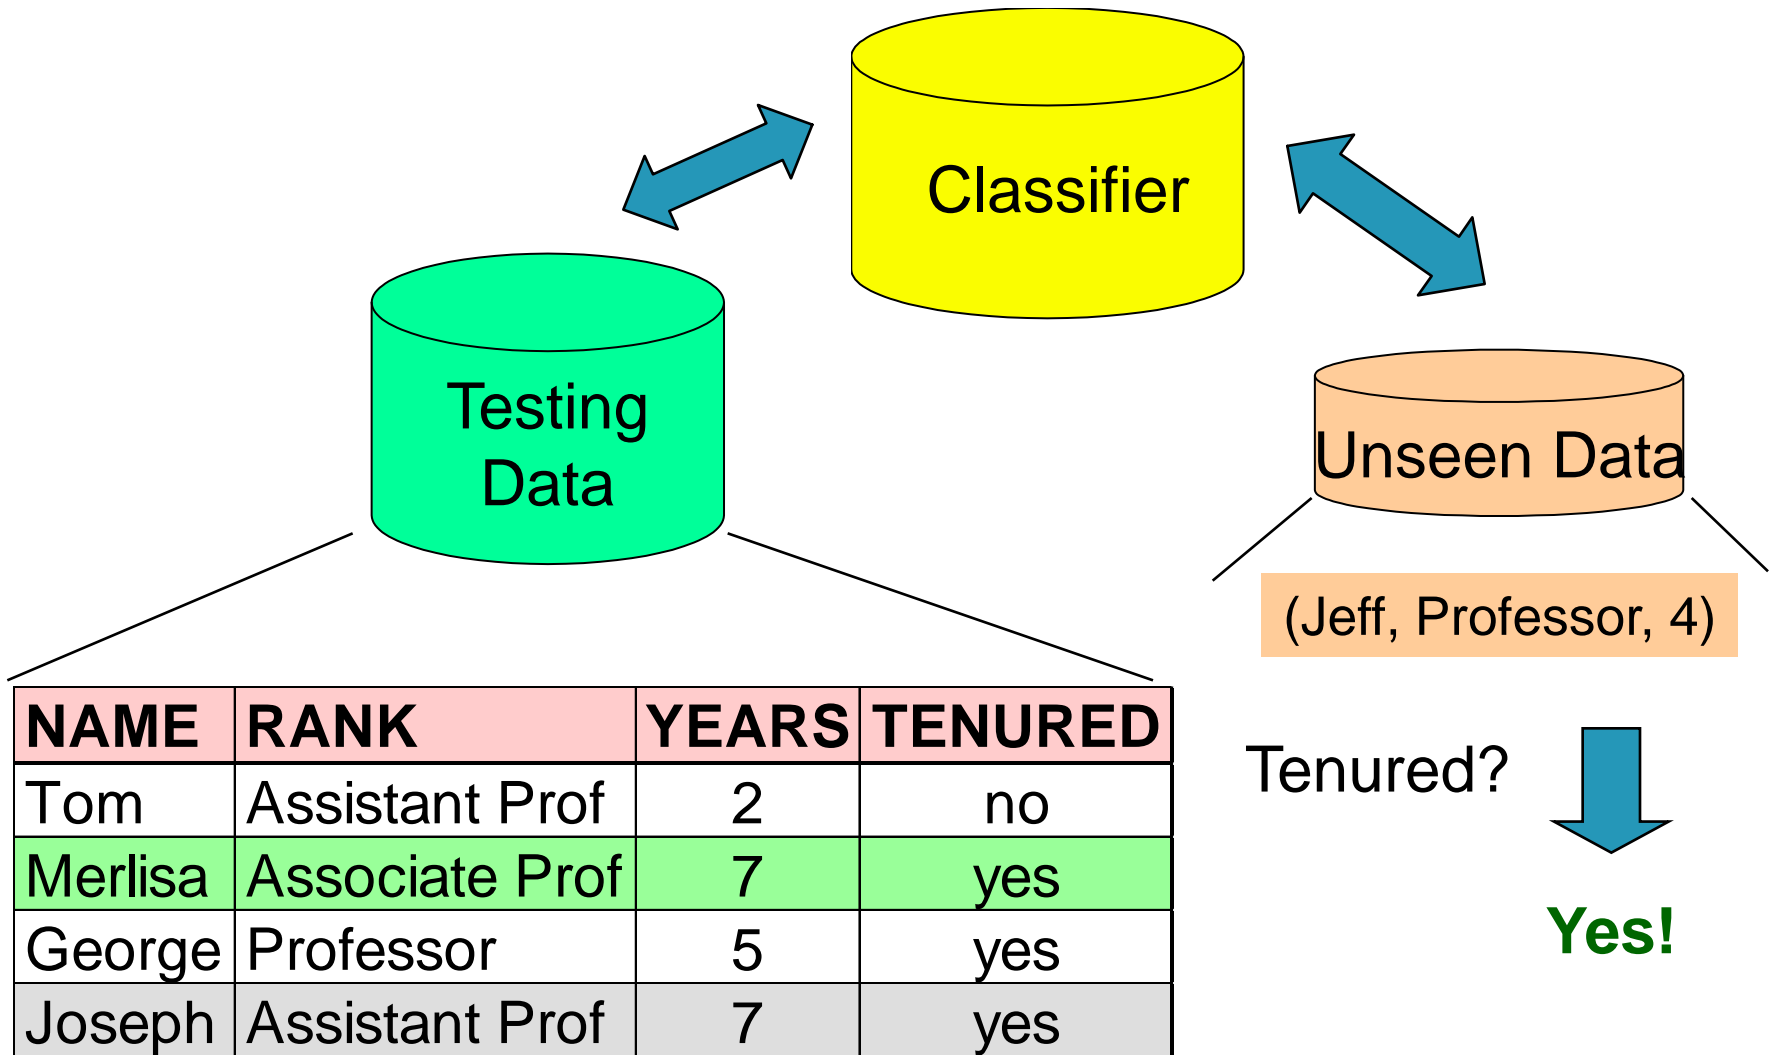
Classification – a Two-Step Process

- **Model construction**: describing a set of predetermined classes
 - Each tuple/sample is assumed to belong to a predefined class, as determined by the **class label attribute**
 - The set of tuples used for model construction is the **training set**
 - The model is represented as decision trees, classification rules, or mathematical formulae (e.g. neural network)
- **Model usage**: for classifying future or unknown objects
 - Evaluate the model
 - The known label of test sample is compared with the classified result from the model
 - E.g., **accuracy** rate is the percentage of test set samples that are correctly classified by the model
 - **Test set** is independent of training set (check for overfitting)
 - If the accuracy is acceptable, use the model to **classify data** tuples whose class labels are not known

Process (1): Model Construction



Process (2): Using the Model



Decision Tree Induction: Training Dataset

*This follows
an example
of Quinlan's
ID3*

age	income	student	credit_rating	buys_computer
<=30	high	no	fair	no
<=30	high	no	excellent	no
31...40	high	no	fair	yes
>40	medium	no	fair	yes
>40	low	yes	fair	yes
>40	low	yes	excellent	no
31...40	low	yes	excellent	yes
<=30	medium	no	fair	no
<=30	low	yes	fair	yes
>40	medium	yes	fair	yes
<=30	medium	yes	excellent	yes
31...40	medium	no	excellent	yes
31...40	high	yes	fair	yes
>40	medium	no	excellent	no

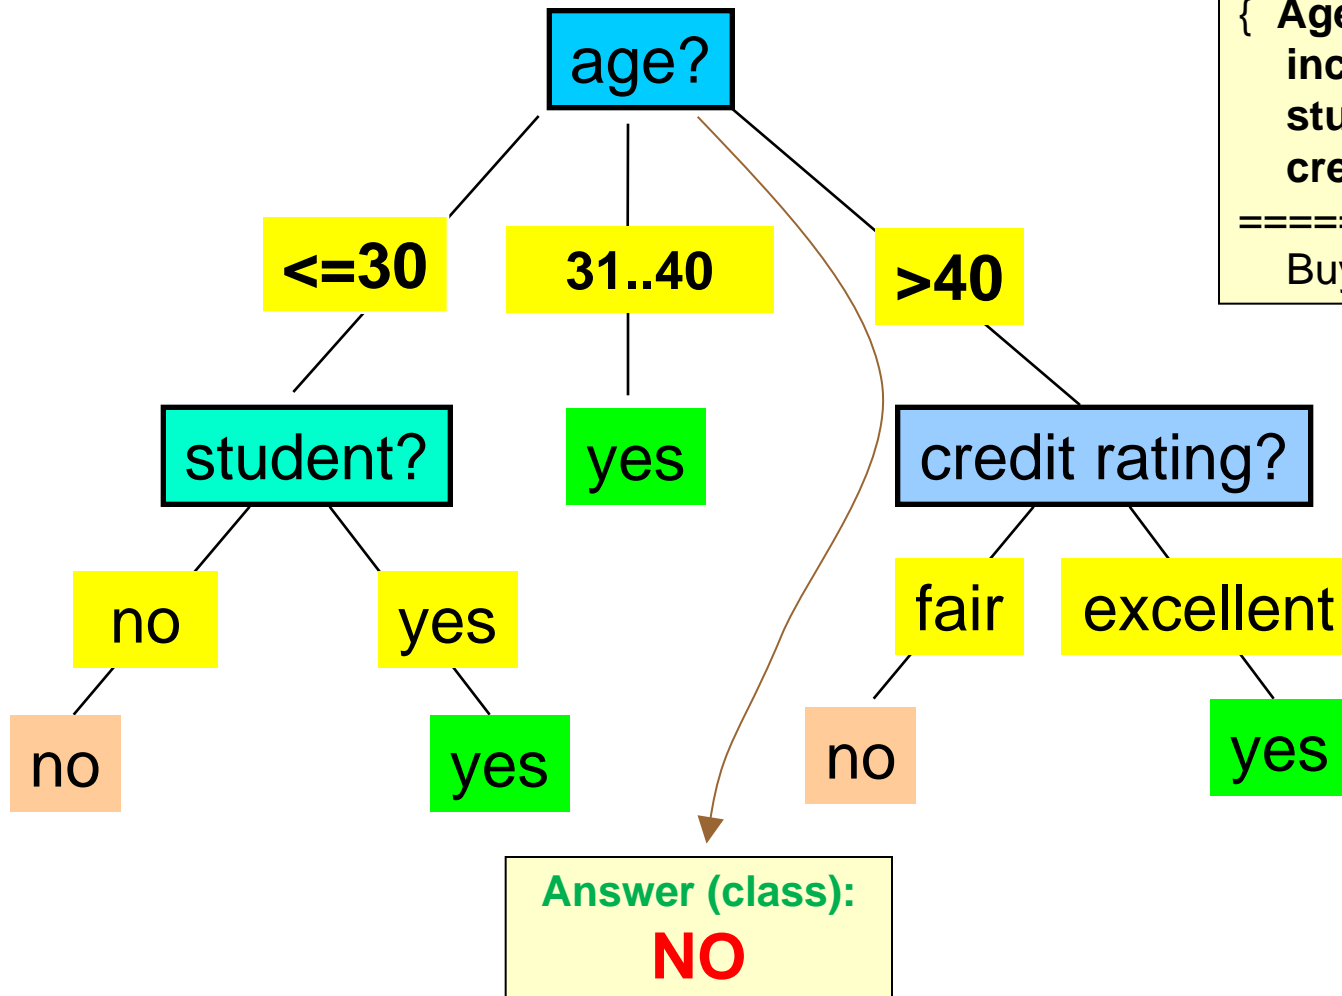
Output: A Decision Tree for “*buys_computer*”

Classify new sample:

{ **Age = 45** and
income = low and
student = no and
credit = fair}

=====

Buy computer = ?



Decision Tree

- Requirements for a Decision Tree algorithm:
 1. Consistent attribute-value description for all data
 2. Predefined classes
 3. Discrete classes
 4. Sufficient data
 5. “Logical” classification models (not weighted decisions)
- Pros
 - **Fast** execution time
 - Generated trees (rules) are **easy to interpret** by humans
 - **Scale well** for large data sets
 - Can handle high dim. data
- Cons
 - Cannot capture **correlations** among attributes
 - Consider only **axis-parallel** cuts

Decision Tree Algorithms

- Classifiers from machine learning and statistical community:
 - ID3
 - C4.5 [Quinlan 93] → C5.0
 - CART (as an advance in applied statistics)
- Classifiers for large databases:
 - SLIQ, SPRINT
 - SONAR
 - Rainforest
- Aspects are **quality** of the tree, **scalability**, and **memory use**.

Algorithm for Decision Tree Induction

- Basic algorithm (a greedy algorithm)
 - Tree is constructed in a **top-down recursive divide-and-conquer manner**
 - At start, all the training examples are at the root
 - Attributes are categorical (if continuous-valued, they are discretized in advance)
 - Examples are partitioned recursively based on selected attributes
 - Test attributes are selected on the basis of a heuristic or statistical measure (e.g., **information gain**)
- Conditions for **stopping partitioning**
 - All samples for a given node belong to the same class
 - There are no remaining attributes for further partitioning
 - **majority voting** is employed for classifying the leaf
 - There are no samples left


Decision Tree Algorithms: First Splitting

education

high school	reject	1
under graduate	accept	2
under graduate	reject	3
graduate	accept	4
graduate	accept	5

salary

10000	reject	1
15000	accept	5
18000	reject	3
40000	accept	2
75000	accept	4



high-school	reject	1
under-graduate	accept	2
under-graduate	reject	3


10000	reject	1
40000	accept	2
18000	reject	3

graduate	accept	4
graduate	accept	5

75000	accept	4
15000	accept	5

we did not explain how we selected “education” attribute for splitting

Reminder... $\log_2 p$



Brief Review of Entropy


- Entropy (Information Theory)
 - **Measure of uncertainty** associated with a random variable
 - Higher entropy \Rightarrow higher uncertainty
 - Lower entropy \Rightarrow lower uncertainty
 - Calculation: For a discrete random variable Y taking m distinct values $\{y_1, \dots, y_m\}$, where $p_i = P(Y=y_i)$

$$H(Y) = \sum_{i=1}^m p_i \cdot \log\left(\frac{1}{p_i}\right) = - \sum_{i=1}^m p_i \cdot \log(p_i)$$

*weighted average
over the classes*

surprise

Brief Review of Entropy



$$H(Y) = \sum_{i=1}^m p_i \cdot \log\left(\frac{1}{p_i}\right) = -\sum_{i=1}^m p_i \cdot \log(p_i)$$

Three examples for $m=2$ classes:

- $p_1=0, p_2=1$

$$H(Y) = -0 \cdot \log(0) - 1 \cdot \log(1) = -0 \cdot (-\infty_{\text{small}}) - 1 \cdot 0 = 0$$

- $p_1=0.25, p_2=0.75$

$$\begin{aligned} H(Y) &= -0.25 \cdot \log(0.25) - 0.75 \cdot \log(0.75) = -0.25 \cdot (-2) - 0.75 \cdot (-0.415) \\ &= 0.811 \end{aligned}$$

- $p_1=0.5, p_2=0.5$

$$H(Y) = -0.5 \cdot \log(0.5) - 0.5 \cdot \log(0.5) = -0.5 \cdot (-1) - 0.5 \cdot (-1) = 1$$

Select the Attribute with highest Information Gain (ID3/C4.5)

- Let p_i be the probability that an arbitrary tuple in D belongs to class C_i , (m classes) estimated by $|C_i \cap D|/|D|$
- Information** (entropy) to classify a tuple **in D** :

$$Info(D) = -\sum_{i=1}^m p_i \log_2(p_i)$$

- Average information** needed, after using attribute A to split D into **k partitions**:

$$Info_A(D) = \sum_{j=1}^k \underbrace{\frac{|D_j|}{|D|}}_{\text{weighted average over the partitions}} \cdot \underbrace{Info(D_j)}_{\text{entropy in partition}}$$

*weighted average
over the partitions*

entropy in partition

- Information gained** by branching on attribute A :

$$Gain(A) = Info(D) - Info_A(D)$$

Information Gain – Example

- Class “buys_computer =yes” (9x)
- Class “buys_computer =no” (5x)

$$Info(D) = I(9,5) = -\frac{9}{14} \log_2\left(\frac{9}{14}\right) - \frac{5}{14} \log_2\left(\frac{5}{14}\right) = 0.94$$

$$Info_{age}(D) = \frac{5}{14} I(2,3) + \underbrace{\frac{4}{14} I(4,0)}_0 + \frac{5}{14} I(3,2)$$

= 0.694

age	yes _i	no _i	I(yes _i , no _i)
<=30	2	3	0,971
31...40	4	0	0
>40	3	2	0,971

“age <=30” has 5 out of 14 samples, with 2 “yes” and 3 “no”

age	income	student	credit_rating	buys_computer
<=30	high	no	fair	no
<=30	high	no	excellent	no
31...40	high	no	fair	yes
>40	medium	no	fair	yes
>40	low	yes	fair	yes
>40	low	yes	excellent	no
31...40	low	yes	excellent	yes
<=30	medium	no	fair	no
<=30	low	yes	fair	yes
>40	medium	yes	fair	yes
<=30	medium	yes	excellent	yes
31...40	medium	no	excellent	yes
31...40	high	yes	fair	yes
>40	medium	no	excellent	no

$$Gain(age) = Info(D) - Info_{age}(D) = \underline{0.246}$$

- Information gains for other splits:

$$Gain(income) = 0.029$$

$$Gain(student) = 0.151$$

$$Gain(credit_rating) = 0.048$$

Decision Tree Algorithm

- Key idea: ***Recursive Partitioning***
 - Take all of your data.
 - Consider *all* possible **values** of *all* **variables**.
 - Select the variable/value ($X=t_1$) that produces the greatest ***separation*** in the target.
 - ($X=t_1$) is called a “split”.
 - If $X < t_1$ then send data point to the “left” branch, otherwise, send data point to the “right” branch.
 - Now repeat same process on these two “nodes”
- You get a “tree”
- Note: CART only uses *binary* splits.


Computing Information-Gain for Continuous-Valued Attributes

- Let attribute A be a continuous-valued attribute
- Must determine the **best split point** for A
 - Sort the value A in increasing order
 - Typically, the **midpoint** between a pair of adjacent values is considered as a possible *split point*
 - $(a_i + a_{i+1})/2$ is the midpoint between the values of a_i and a_{i+1}
 - The point with the *minimum expected information requirement* for A is selected as the split-point for A
- Split:
 - D1 is the set of tuples in D satisfying $A \leq \text{split-point}$, and D2 is the set of tuples in D satisfying $A > \text{split-point}$

Attribute Selection Measure Comparison

- **Information gain** (ID3/C4.5)
 - All attributes are assumed to be categorical.
 - Can be modified for continuous-valued attributes.
- **Gini impurity** (IBM Intelligent Miner, CART, SLIQ, SPRINT)
 - All attributes are assumed continuous-valued.
 - Can be modified for categorical attributes.
 - Assume there exist several possible split values for each attribute.

Gini



- Corrado Gini, Italian statistician, 1884-1965
- ***Gini coefficient***
 - Used to show inequality of income distribution in a population
 - Large if unequal incomes,
small if equal incomes
- ***Gini impurity***
 - A measure for the distribution of labels in a set
 - Large if many equally distributed labels,
small if large probability only for few labels
- Hence, Gini index \neq Gini coefficient
 - **Attention:** Both sometimes referred to as “Gini index”

Attribute Selection using *Gini Impurity*

- A data set D contains examples from m classes, and
- p_j is the relative frequency of class j in D .
- Then Gini impurity, $Gini(D)$, is defined as:

$$Gini(D) = \underbrace{\sum_{j=1}^m p_j}_{\text{weighted average over the classes}} \underbrace{(1 - p_j)}_{\text{probability of incorrect labeling}} = 1 - \sum_{j=1}^m p_j^2$$


- Gini measures how often a randomly chosen element from the set would be incorrectly labeled if it was randomly labeled according to the distribution of labels in the subset.
- Should be minimized!
- Note: while $\sum_i p_i = 1$ always, not so $\sum_i p_i^2$

Attribute Selection using *Gini Impurity*

$$Gini(D) = \sum_{j=1}^m p_j(1 - p_j) = 1 - \sum_{j=1}^m p_j^2$$

- **Minimum** ($1-1=0$) when all records belong to one class
→ **most interesting** information for a split
- **Maximum** ($1 - 1/m$) when records are *equally distributed* among all classes
→ **least interesting** information

Gini impurity for $m=2$ classes



Splitting Based on *Gini* Impurity

- When a node p is split into k partitions, the *quality* of split is computed as

$$Gini_{\text{split}} = \sum_{i=1}^k \frac{n_i}{n} Gini(i)$$

where: n_i = number of records at child i ,
 n = number of records at node p .

- Interpretation: weighted sum of *Gini* impurity for subsets i of samples caused by splitting
- If a data set D is split into two subsets D_1 and D_2 with sizes n_1 and n_2 respectively, the *Gini* impurity $Gini(D)$ is defined as

$$Gini_{\text{split}}(D) = \frac{n_1}{n} Gini(D_1) + \frac{n_2}{n} Gini(D_2)$$

Splitting Based on *Gini* Impurity

- Need to enumerate all possible splitting points for each attribute
- The attribute that provides the smallest $Gini_{split}(D)$ is chosen to split the node

Gini Impurity – Example

- Class “buys_computer =yes” (9x)
- Class “buys_computer =no” (5x)

$$Gini(D) = Gini(9,5) = 1 - \left(\frac{9}{14}\right)^2 - \left(\frac{5}{14}\right)^2 = 0.459$$

age	yes _i	no _i	Gini(yes _i ,no _i)
<=30	2	3	0,48
31...40	4	0	0
>40	3	2	0,48

$$Gini_{age}(D) = \frac{5}{14} Gini(2,3) + \frac{4}{14} \underbrace{Gini(4,0)}_0 + \frac{5}{14} Gini(3,2) = 0.343$$

“age <=30” has 5 out of 14 samples, with 2 “yes” and 3 “no”

age	income	student	credit_rating	buys_computer
<=30	high	no	fair	no
<=30	high	no	excellent	no
31...40	high	no	fair	yes
>40	medium	no	fair	yes
>40	low	yes	fair	yes
>40	low	yes	excellent	no
31...40	low	yes	excellent	yes
<=30	medium	no	fair	no
<=30	low	yes	fair	yes
>40	medium	yes	fair	yes
<=30	medium	yes	excellent	yes
31...40	medium	no	excellent	yes
31...40	high	yes	fair	yes
>40	medium	no	excellent	no

- Compute Gini for other splits:

$$Gini(income) = \dots$$

$$Gini(student) = \dots$$

$$Gini(credit_rating) = \dots$$

- Consider also other splits, e.g. age {<=30 & 31...40} and {>40}
- Split at lowest value

Examples for Computing Gini (for $m=2$ classes)

- Gini Impurity for a given node t : $Gini(t) = 1 - \sum_{j=1}^m p(j | t)^2$

$p(j | t)$ is the *relative frequency* of class j at node t

C1	0
C2	6

$$P(C1) = 0/6 = 0 \quad P(C2) = 6/6 = 1$$

$$\mathbf{Gini} = 1 - P(C1)^2 - P(C2)^2 = 1 - 0 - 1 = \mathbf{0}$$

Minimum

C1	1
C2	5

$$P(C1) = 1/6 \quad P(C2) = 5/6$$

$$\mathbf{Gini} = 1 - (1/6)^2 - (5/6)^2 = \mathbf{0.278}$$

C1	2
C2	4

$$P(C1) = 2/6 \quad P(C2) = 4/6$$

$$\mathbf{Gini} = 1 - (2/6)^2 - (4/6)^2 = \mathbf{0.444}$$

C1	3
C2	3

$$P(C1) = 3/6 \quad P(C2) = 3/6$$

$$\mathbf{Gini} = 1 - (3/6)^2 - (3/6)^2 = \mathbf{0.5}$$

Maximum


Making decisions – Errors in value

Fresher after lunch




Dan Gilbert: Why we make bad decisions,
TED talks [Video online](#)

Search



Matlab Demo of Decision Tree

- In Matlab, `t = classregtree(X,y,'Name',value)` creates a decision tree.
- **Example:** Create a classification tree for Fisher's iris data, a typical test case for many classification techniques.




Matlab Demo of Decision Tree

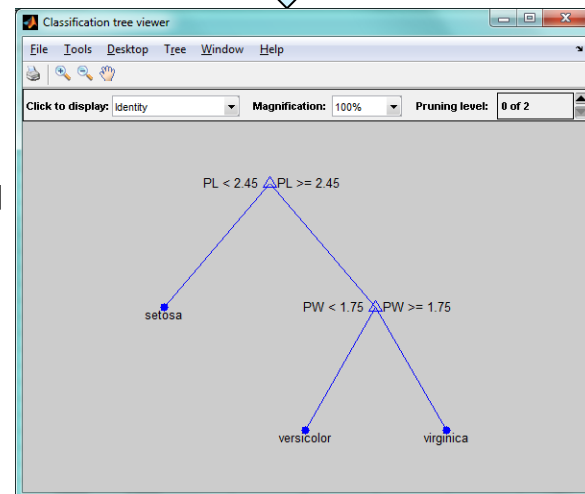
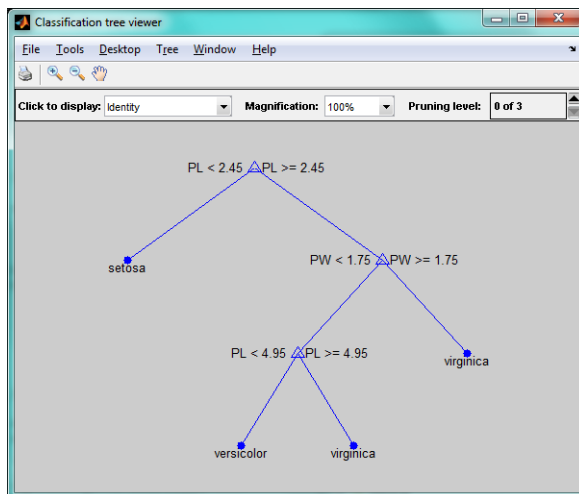
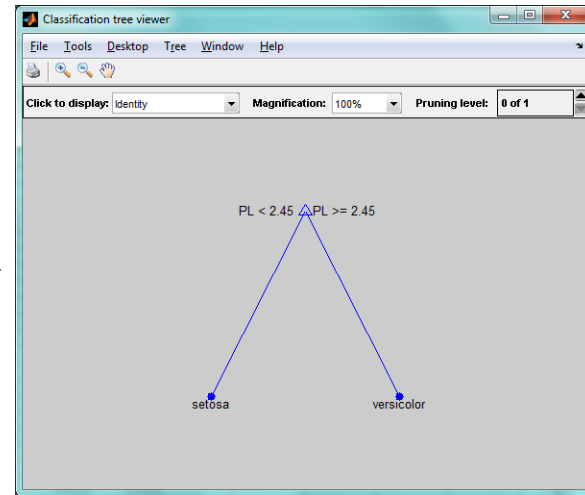

- In Matlab, `t = classregtree(X,y,'Name',value)` creates a decision tree.
- **Example:** Create a classification tree for Fisher's iris data, a typical test case for many classification techniques.
 - In this data set, four attributes (Sepal Length, Sepal Width, Petal Length and Petal Width) are considered in order to distinguish three species of flowers (*Iris setosa*, *Iris virginica* and *Iris versicolor*).
 - Commands:

```
load fisheriris;  
t = classregtree(meas,species,... 'names',{ 'SL'  
      'SW' 'PL' 'PW' } );
```
 - Program generates a decision tree based on the data set.

Matlab Demo of Decision Tree



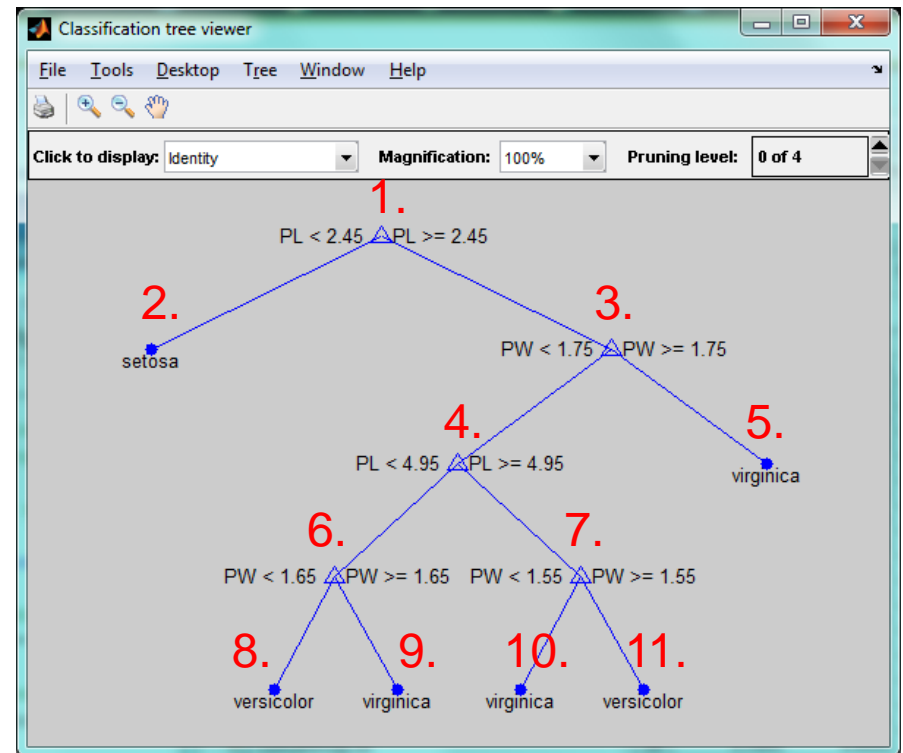
Matlab Demo of Decision Tree



Matlab Demo of Decision Tree


Final Decision tree for classification

1. if $PL < 2.45$ then node 2 elseif $PL \geq 2.45$ then node 3
2. class = setosa
3. if $PW < 1.75$ then node 4 elseif $PW \geq 1.75$ then node 5
4. if $PL < 4.95$ then node 6 elseif $PL \geq 4.95$ then node 7
5. class = virginica
6. if $PW < 1.65$ then node 8 elseif $PW \geq 1.65$ then node 9
7. if $PW < 1.55$ then node 10 elseif $PW \geq 1.55$ then node 11
8. class = versicolor
9. class = virginica
10. class = virginica
11. class = versicolor



Matlab Demo of Decision Tree

- We can also prune the tree to avoid overfitting
- `tt = prune(t,'level',2)`



Advanced Decision Trees


- C4.5
 - Improved handling of continuous variables
 - C source code available
- C5
 - Quinlan made further improvements (boosting)
 - Many commercial data mining packages use the C5 algorithm
 - Source code available at a cost!
- CART
 - Breiman et al (Classification & regression trees, 1984)
 - similar to C4.5, boosting & bagging the data sets

Decision Tree Algorithms – C4.5


- **Recursive building** tree phase:
 1. *Initialize root node of tree.*
 2. **while** a node *N* that can be split:
 3. **for each** attribute *A*, evaluate splits on *A*,
 4. use best split to split *N*.
- Use **entropy (information gain)** to find best split
- Separate attribute lists maintained in each node of tree

C4.5 – Possible Mechanisms for Tests


- a. “standard” test on a **discrete attribute**:
one branch for each possible value of that attribute



- b. If attribute Y has **continuous numeric values**,
binary test with outcomes $Y \leq Z$ and $Y > Z$ could
be defined



- c. possible values are allocated to a variable
number of **groups** with one outcome/branch
for each group



New example (1) Threshold Finding with Gain

Sometimes we have to find the threshold and the attribute

Database D

Attribute 1	Attribute 2	Attribute 3	Class
A	70	True	Class1
A	90	True	Class2
A	85	False	Class2
A	95	False	Class2
A	70	False	Class1
B	90	True	Class1
B	78	False	Class1
B	65	True	Class1
B	75	False	Class1
C	80	True	Class2
C	70	True	Class2
C	80	False	Class1
C	80	False	Class1
C	96	False	Class1

Attribute 2:

- After a sorting process, the set of values is: {65, 70, 75, 78, 80, 85, 90, 95, 96},
- ... the set of potential threshold values Z is: {65, 70, 75, 78, 80, 85, 90, 95}.
- The optimal Z value is $Z=80$ (highest Inf. Gain)

9 are ≤ 80 ... 7 are Class1 2 are Class2

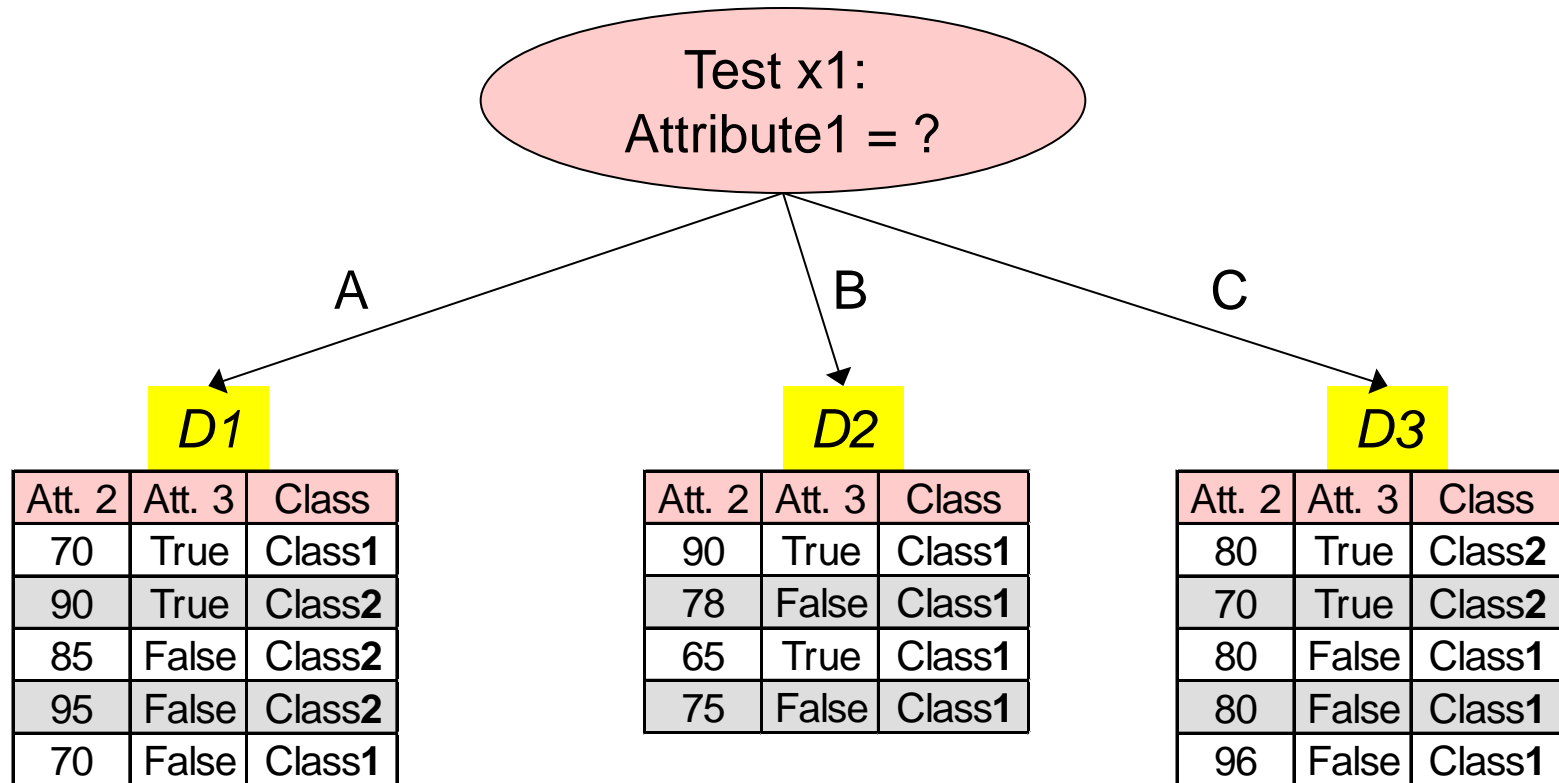
5 are > 80

- $\text{Info}_{Z=80}(D) = \frac{9}{14} \cdot (-\frac{7}{9} \cdot \log_2(\frac{7}{9}) - \frac{2}{9} \cdot \log_2(\frac{2}{9})) + \frac{5}{14} \cdot (-\frac{2}{5} \cdot \log_2(\frac{2}{5}) - \frac{3}{5} \cdot \log_2(\frac{3}{5})) = 0.837$ bits
- $\text{Gain}(Z=80) = 0.940 - 0.837 = 0.103$ bits

However, Attribute 1 gives the highest gain of 0.246 bits, and therefore this attribute will be selected for the first splitting

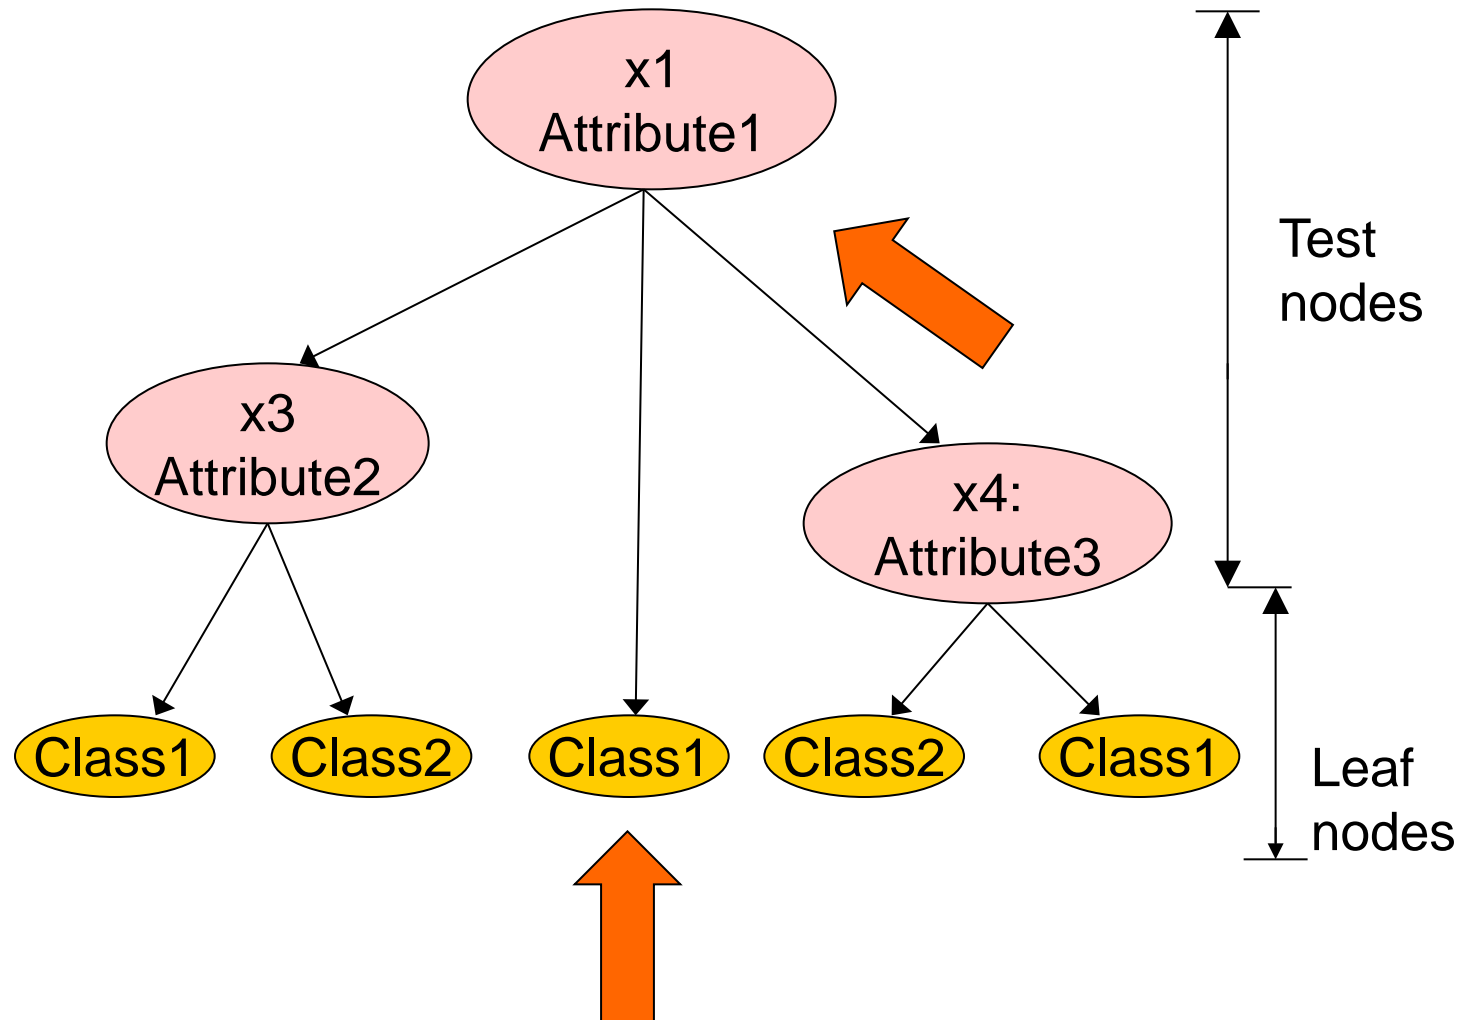
(are attributes with many values favoured in general ... ?)

New Example (2) Initial Decision Tree



Initial decision tree and subset cases for a database D

New example (3) Final Decision Tree



All of them are in CLASS1

Final Decision Tree as Pseudo Code

- Decision Tree – **Pseudo-code Example:**

```
If      Attribute1 = A
  Then
    If      Attribute2 <= 70
      Then
        Classification = CLASS1;
      Else
        Classification = CLASS2;
    Elseif  Attribute1 = B
      Then
        Classification = CLASS1;
    Elseif  Attribute1 = C
      Then
        If      Attribute3 = True
          Then
            Classification = CLASS2;
          Else
            Classification = CLASS1.
```

C4.5 Algorithm: Gain Ratio

- *Revision:* Measures we defined so far:

- Entropy to classify a tuple in D :
- Information needed (after using A to split D into k partitions) to classify D :
- Information gained for attribute A :

$$Info(D) = -\sum_{i=1}^m p_i \log_2(p_i)$$


$$Info_A(D) = \sum_{j=1}^k \frac{|D_j|}{|D|} \cdot Info(D_j)$$

$$Gain(A) = Info(D) - Info_A(D)$$

- Information gain measure is **biased** towards attributes with a large number of values (*this is also true for the Gini index*)
- C4.5 (a successor of ID3) uses gain ratio to **normalize the information gain**

$$SplitInfo = -\sum_{j=1}^k \left(\frac{|D_j|}{|D|} \cdot \log_2 \left(\frac{|D_j|}{|D|} \right) \right)$$

$$GainRatio(A) = Gain(A) / SplitInfo(A)$$



Information Gain → Gain Ratio (prev. Example)

- Class “buys_computer =yes” (9x)
- Class “buys_computer =no” (5x)

$$Info(D) = I(9,5) = -\frac{9}{14} \log_2\left(\frac{9}{14}\right) - \frac{5}{14} \log_2\left(\frac{5}{14}\right) = 0.94$$

$$Info_{age}(D) = \frac{5}{14} I(2,3) + \frac{4}{14} I(4,0) + \frac{5}{14} I(3,2) = 0.694$$

age	yes _i	no _i	I(yes _i , no _i)
<=30	2	3	0,971
31...40	4	0	0
>40	3	2	0,971

“age <=30” has 5 out of 14 samples, with 2 “yes” and 3 “no”

age	income	student	credit_rating	buys_computer
<=30	high	no	fair	no
<=30	high	no	excellent	no
31...40	high	no	fair	yes
>40	medium	no	fair	yes
>40	low	yes	fair	yes
>40	low	yes	excellent	no
31...40	low	yes	excellent	yes
<=30	medium	no	fair	no
<=30	low	yes	fair	yes
>40	medium	yes	fair	yes
<=30	medium	yes	excellent	yes
31...40	medium	no	excellent	yes
31...40	high	yes	fair	yes
>40	medium	no	excellent	no

$$Gain(age) = Info(D) - Info_{age}(D) = 0.246$$

$$SplitInfo(age) = -\sum_{j=1}^3 \left(\frac{|D_j|}{|D|} \cdot \log_2 \left(\frac{|D_j|}{|D|} \right) \right) = -\frac{5}{14} \log_2\left(\frac{5}{14}\right) - \frac{4}{14} \log_2\left(\frac{4}{14}\right) - \frac{5}{14} \log_2\left(\frac{5}{14}\right) = 1.577$$

$$GainRatio(age) = 0.246 / 1.557 = 0.156$$

C4.5 Algorithm for Continuous Numeric Values

- Define binary test with outcomes $X \leq Z$ and $X > Z$, based on comparing the value of attribute against a **threshold value** Z
- **Sort** the training samples w.r.t. the values of the chosen attribute X
 - Number of these values is finite
 - Notation for sorted order: $\{v_1, v_2, \dots, v_m\}$
- Examine **all splits** to find the **optimal split**
 - $m-1$ possible splits on X .
 - Any threshold value between v_i and v_{i+1} has the same effect of dividing the cases into $\{v_1, v_2, \dots, v_i\}$ and $\{v_{i+1}, v_{i+2}, \dots, v_m\}$.
- Normal choice as representative threshold:
midpoint of each interval: $(v_i + v_{i+1})/2$
 - C4.5 chooses the **smaller** value v_i of an interval $\{v_i, v_{i+1}\}$, rather than the midpoint – ensures that threshold values exist in the data


C4.5 Algorithm: Unknown Values

- In C4.5, samples with unknown values are ***distributed probabilistically*** according to the ***relative frequency*** of known values

- New information gain criterion for split in attribute X :

$$Gain(X) = F \cdot (Info(D) - Info_X(D))$$

- ***Factor F*** = number of samples in database with known value for a given attribute / total number of samples in a data set
- ***Factor F*** here 13/14



Attribute 1	Attribute 2	Attribute 3	Class
A	70	True	Class1
A	90	True	Class2
A	85	False	Class2
A	95	False	Class2
A	70	False	Class1
?	90	True	Class1
B	78	False	Class1
B	65	True	Class1
B	75	False	Class1
C	80	True	Class2
C	70	True	Class2
C	80	False	Class1
C	80	False	Class1
C	96	False	Class1

C4.5 Algorithm: Unknown Values – Example (1)

13 remaining cases with values for Attribute1



$$\text{Info}(D) = -8/13 \log_2 (8/13) - 5/13 \log_2 (5/13) = \mathbf{0.961 \text{ bits}}$$



8 belong to CLASS1



5 belong to CLASS2

Test X_1 for the three values A, B, or C:

$$\begin{aligned} \text{Info}_{X_1}(D) &= 5/13 (-2/5 \log_2 (2/5) - 3/5 \log_2 (3/5)) \\ &\quad + 3/13 (-3/3 \log_2 (3/3) - 0/3 \log_2 (0/3)) \\ &\quad + 5/13 (-3/5 \log_2 (3/5) - 2/5 \log_2 (2/5)) \\ &= \mathbf{0.747 \text{ bits}} \end{aligned}$$

$$\mathbf{\text{Gain}(X_1) = 13/14 \cdot (0.961 - 0.747) = 0.199 \text{ bits}}$$



Factor F

Attribute 1	Attribute 2	Attribute 3	Class
A	70	True	Class1
A	90	True	Class2
A	85	False	Class2
A	95	False	Class2
A	70	False	Class1
?	90	True	Class1
B	78	False	Class1
B	65	True	Class1
B	75	False	Class1
C	80	True	Class2
C	70	True	Class2
C	80	False	Class1
C	80	False	Class1
C	96	False	Class1

C4.5 Algorithm: Unknown Values – Example (2)

Attribute 1	Attribute 2	Attribute 3	Class
A	70	True	Class1
A	90	True	Class2
A	85	False	Class2
A	95	False	Class2
A	70	False	Class1
?	90	True	Class1
B	78	False	Class1
B	65	True	Class1
B	75	False	Class1
C	80	True	Class2
C	70	True	Class2
C	80	False	Class1
C	80	False	Class1
C	96	False	Class1

Distribution of samples into subsets with corresponding weight factors w

D1: Attribute1 = A

Att.2	Att.3	Class	w
70	True	Class1	1
90	True	Class2	1
85	False	Class2	1
95	False	Class2	1
70	False	Class1	1
90	True	Class1	5/13

D1: Attribute1 = B

Att.2	Att.3	Class	w
90	True	Class1	3/13
78	False	Class1	1
65	True	Class1	1
75	False	Class1	1

D1: Attribute1 = C

Att.2	Att.3	Class	w
80	True	Class2	1
70	True	Class2	1
80	False	Class1	1
80	False	Class1	1
96	False	Class1	1
90	True	Class1	5/13

C4.5 Algorithm: Generalizing Partitioning

- When a sample from **D with known value** is assigned to subset D_i , its probability belonging to **D_i is 1**, and in all other subsets is 0
- C4.5 associates with each sample (having **missing value**) a **weight w** representing the **probability** that it belongs to each subset D_i :

$$w_{\text{new}} = w_{\text{old}} \cdot P(D_i)$$

- Splitting set D using test X_1 on Attribute1: New weights w_i will be probabilities, here: 5/13, 3/13, and 5/13, since initial w_{old} is 1

$$|D_1| = 5 + 5/13, \quad |D_2| = 3 + 3/13, \quad \text{and} \quad |D_3| = 5 + 5/13$$

- The decision tree **leaves** are defined with two new parameters: **$(|D_i|/E)$**
- **$|D_i|$** is the sum of the **fractional samples** that reach the leaf, and **E** is the **number of samples** belonging to classes other than nominated class
- (3.4 / 0.4) means:
 - 3.4 (or $3 + 5/13$) fractional training samples reached leaf,
 - 0.4 (or $5/13$) of which did not belong to the class of the leaf

Partitioning – Example

- Decision tree for the database D with missing values:

```
If      Attribute1 == A
  Then
    If      Attribute2 <= 70
      Then
        Classification = CLASS1      (2.0 / 0);
      Else
        Classification = CLASS2      (3.4 / 0.4);
    Elseif Attribute1 == B
      Then
        Classification = CLASS1      (3.2 / 0);
    Elseif Attribute1 == C
      Then
        If      Attribute3 = True
          Then
            Classification = CLASS2    (2.4 / 0.4);
          Else
            Classification = CLASS1    (3.0 / 0).
```

($|D_i|/E$):

$|D_i|$ = sum of the fractional samples that reach the leaf,

E = number of samples that belong to classes other than the nominated class.

Enhancements to Basic Decision Tree Induction (Intermediate Summary)

- Allow for ***continuous-valued attributes***
 - Partition the continuous attribute value into a discrete set of intervals, dynamically defined using the data's attribute values
- Handle ***missing attribute values***
 - Assign probability to each of the possible values
- ***Attribute construction***
 - Create new attributes based on existing ones that are sparsely represented
 - This reduces fragmentation, repetition, and replication
- Challenge: ***incremental learning*** of decision trees

Decision Tree Algorithms – Building and Pruning

- ***Building phase***

- Recursively split nodes using best splitting attribute for node.

- ***Pruning phase***

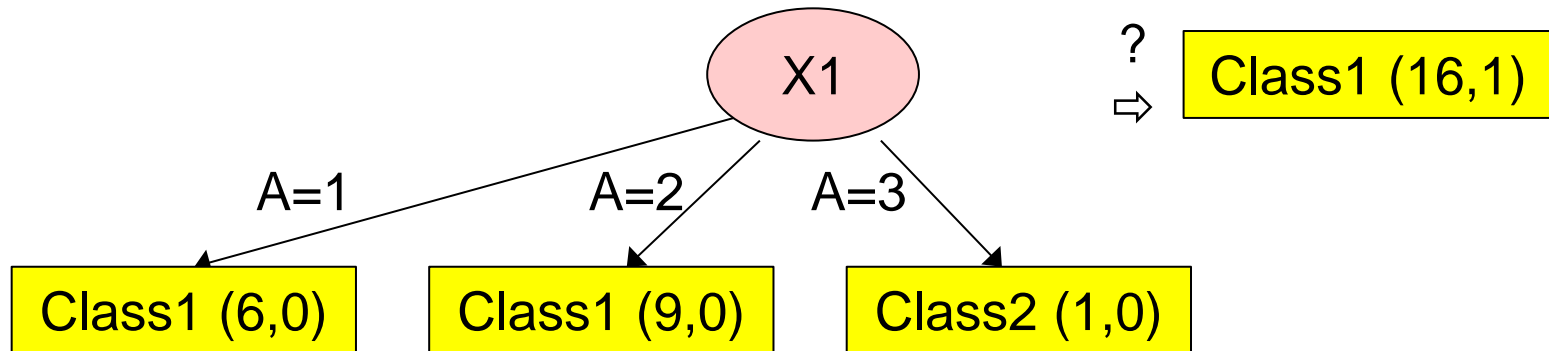
- Smaller imperfect decision tree generally achieves better accuracy on test data.
- Prune leaf nodes recursively to prevent over-fitting.

Avoid Overfitting in Classification

- The generated tree may overfit the training data:
 - Too **many branches**, some may reflect anomalies due to noise or outliers
 - Result: poor accuracy for unseen samples
- Two approaches to avoid overfitting:
 - **Prepruning**: Halt tree construction early—do not split a node if this would result in the goodness measure falling below a threshold
 - Difficult to choose an appropriate threshold
 - **Postpruning**: Remove branches from a “fully grown” tree—get a sequence of progressively pruned trees
 - Use a set of data different from the training data to decide which is the “best pruned tree”

Pruning a Decision Tree

- **Pruning**: Discarding one or more subtrees and replacing them with leaves
 - C4.5 follows a **postpruning** approach (pessimistic pruning)




will we replace this subtree with a single leaf node?

Pruning Decision Tree: Predicted Error

$$PE = \sum_{i=1}^{nodes} n_i \cdot U_{25\%}$$

of samples in the node

upper limit on error rate (for the node):
from statistical tables for binomial distributions



- Using default confidence of 25%, **upper limits on the error rates** for all nodes are collected from statistical tables for binomial distributions:

Tree: $U_{25\%}(6,0) = 0.206$, $U_{25\%}(9,0) = 0.143$, $U_{25\%}(1,0) = 0.750$


Node: $U_{25\%}(16,1) = 0.157$

- Predicted errors** for the subtree and the replaced node are:
 - $PE_{tree} = 6 \cdot 0.206 + 9 \cdot 0.143 + 1 \cdot 0.750 = 3.257$
 - $PE_{node} = 16 \cdot 0.157 = 2.512$
 - Since $PE_{tree} > PE_{node}$, replace the subtree with the new leaf node.

$$U_{CF}(|D_i|, E)$$

- Consider classifying E examples incorrectly out of $|D_i|$ samples (like observing E events in $|D_i|$ trials in the binomial distribution)
- For a given confidence level CF, the upper limit on the error rate over the whole population is $U_{CF}(|D_i|, E)$ with CF% confidence.
- Example:

- $U_{25\%}(100, 6)$
- 100 examples in a leaf
- 6 examples misclassified
- How large is the true error assuming a pessimistic estimate with a confidence of 25%?



Extracting Decision Rules from Trees

- Represent the knowledge in the form of **IF-THEN** rules
 - One rule is created for each **path from the root to a leaf**.
 - Each **attribute-value pair** along a path forms a **conjunction**.
 - The leaf node holds the class prediction.
- **Rules are easier** for humans to understand

Examples:

```
IF age = "<=30" AND student = "no"  
THEN buys_computer = "no"
```

```
IF age = "<=30" AND student = "yes"  
THEN buys_computer = "yes"
```

```
IF age = "31...40"  
THEN buys_computer = "yes"
```

```
IF age = ">40" AND credit_rating = "excellent"  
THEN buys_computer = "yes"
```

```
IF age = ">40" AND credit_rating = "fair"  
THEN buys_computer = "no"
```

Rule Ordering


If more than one rule is triggered, we need **conflict resolution**

- Size ordering
 - assign the highest priority to the triggering rules that has the “toughest” requirement (i.e., with the *most attribute tests*)
- Class-based ordering
 - Rules for the most frequent class come first, or
 - Sort based on misclassification cost per class
- Rule-based ordering (decision list)
 - rules are organized into one long priority list, according to some measure of rule quality (e.g. accuracy, # attribute tests) or by experts

C4.5 Algorithm: Generating Decision Rules

may not really simplify

Decision tree



Transformation
 Paths into Rules

Decision rules

If A=1 and B=1	Then Class1
If A=1 and B=2	Then Class2
If A=2	Then Class1



Decision rules
for database D:

Attribute 1	Attribute 2	Attribute 3	Class
A	70	True	Class1
A	90	True	Class2
A	85	False	Class2
A	95	False	Class2
A	70	False	Class1
?	90	True	Class1
B	78	False	Class1
B	65	True	Class1
B	75	False	Class1
C	80	True	Class2
C	70	True	Class2
C	80	False	Class1
C	80	False	Class1
C	96	False	Class1

<i>If</i>	Attribute1 = A and Attribute2 ≤ 70	<i>Then</i>	Classification = CLASS1 (2.0 / 0);
<i>If</i>	Attribute1 = A and Attribute2 > 70	<i>Then</i>	Classification = CLASS2 (3.4 / 0.4);
<i>If</i>	Attribute1 = B	<i>Then</i>	Classification = CLASS1 (3.2 / 0);
<i>If</i>	Attribute1 = C and Attribute3 = True	<i>Then</i>	Classification = CLASS2 (2.4 / 0.4);
<i>If</i>	Attribute1 = C and Attribute3 = False	<i>Then</i>	Classification = CLASS1 (3.0 / 0).

Bottom example is for previous partitioning data set (14 samples).⁷⁴


Limitations of Decision Trees and Decision Rules (1)



Example:

- 2D samples are classified using a third dimension for classes
- Problematic: classification function is much **more complex** with **related attributes**

Limitations of Decision Trees and Decision Rules (2)



Limitations of Decision Trees and Decision Rules (3)

- Let a given class be supported, if **any k out of n** conditions are met.
- To represent this classifier with rules, it would be necessary to define $\binom{n}{k}$ regions only for one class

$$\binom{n}{k} = \frac{n!}{k! (n - k)!}$$

- **Example:** Medical diagnostic:
 - If 4 out of 11 symptoms support diagnosis of a given disease, then the corresponding classifier will generate 330 regions in 11-dimensional space for positive diagnosis only.
- ⇒ corresponds to 330 decision rules.

Limitations of Decision Trees and Decision Rules: Further Ideas

- Introducing new attributes, rather than removing old ones, can avoid sometimes-intensive fragmentation of the n-dimensional space:

Model: $(A1 \vee A2 \vee A3) \wedge (A4 \vee A5 \vee A6) \wedge (A7 \vee A8 \vee A9) \rightarrow \mathbf{C1}$

Solution 1:

$A1 \wedge A4 \wedge A7 \rightarrow C1$	}	\rightarrow	27 combinations
$A1 \wedge A5 \wedge A7 \rightarrow C1$			
$A1 \wedge A6 \wedge A7 \rightarrow C1$			
...			

Solution 2: Introduce **new derived attributes**:

$B1 = A1 \vee A2 \vee A3$

$B2 = A4 \vee A5 \vee A6$

$B3 = A7 \vee A8 \vee A9$

$\rightarrow \mathbf{B1 \wedge B2 \wedge B3 \rightarrow C1}$

Decision Trees (Summary)


■ Advantages


- Automatically create tree representations from data
- Trees can be converted to rules, can discover “new” rules
- Identify most discriminating attribute first
 - Using Information Gain (Ratio) or Gini Impurity
- Tree can handle discrete, continuous, mixed, and missing attributes

■ Disadvantages

- Trees can become large and difficult to understand
- Can produce counter-intuitive rules
- Examines attributes individually, but not inter-attribute relationships
- Future splits not known when splitting → not globally optimal tree
- Tree induction rules have no direct relation to training objective, i.e. minimizing the classification error

Limitations: Decisions over Time

 Errors in Value



COMPARING WITH
the **Past** instead of the **Possible**

Dan Gilbert: Why we make bad decisions,
TED talks, 2008. [Video online](#)