


# Data Mining

## Lecture 5

### Classification with Supervised Neural Networks



<http://www.informatik.uni-hamburg.de/WTM/>

# Why Learning? Some Quotes


- “Artificial Intelligence is realised only when a computer can ‘discover’ *for itself* new techniques for problem solving” Fogel (1966)
- “Intelligent agents must be able to *change* through the course of their interactions with the world” Luger (2002)
- “A machine or software tool would not be viewed as intelligent if it could not *adapt to* changes in its environment” Callan (2003)

# What is Neural Learning?


- Modify and improve behaviour by past experience
- How does the brain learn?
  - Strengths of synaptic connections vary
- Hebb's rule
  - If two neurons connected by a synapse fire simultaneously then the synapse strengthens
  - If two neurons connected by a synapse do not fire simultaneously then the synapse weakens
  - “fire together, wire together”

# Learning Regression Problems

- Curve Fitting (with *noise*)
- Function Approximation
- Many other functions could fit the data





# Learning Classification Problems



# Computer versus Brain


- The ***von Neumann architecture*** uses a single processing unit
  - Floating Point Operations Per Second (typical today: 1 Tera FLOPS,  $10^{12}$ )
  - Absolute arithmetic precision




## The ***brain***

- Uses many but slow, unreliable processors acting in parallel but they produce robust learned behaviour

# A Real Neuron




# 3D-View of Neurons in the Brain



Shen, H. See-through brains clarify connections. *Nature*, vol. 496, pp. 151, Macmillan Publishers Limited, 11 April 2013. [Video online](#)




# The Neuron



# A Neuron's Firing



# Noninvasive Inspections of the Brain



# Parallel Processing in the Brain

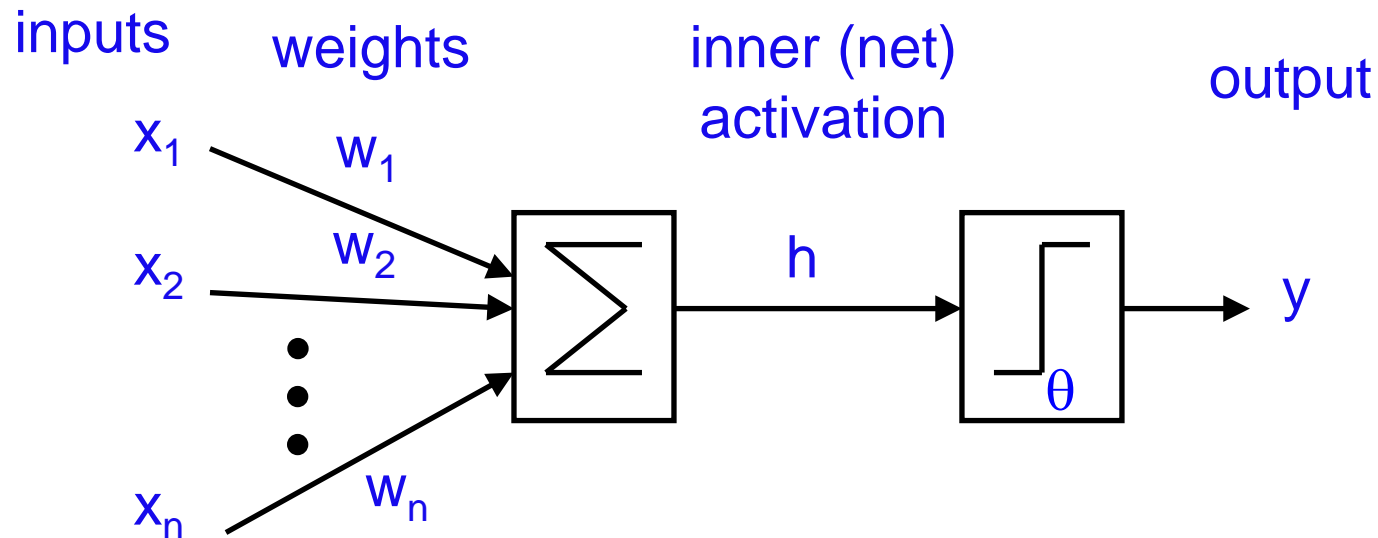
The human brain:

- Weight on average 1.4kg
- Contains around  **$10^{11}$  neurons**
  - Many ***different types***
  - In computational terms,  $10^{11}$  simple processors
    - Each takes a few milliseconds to do a computation
    - But the whole brain is very fast
- Has about  **$10^{14}$  synapses**
- ***Highly connected***
  - Things done massively in parallel
  - Robust to faults

# Neuron Activity

## Neuron Activity The Synapse

# Perceptron Neurons



Greatly simplified biological neurons

- Sum the inputs  $x_j$  each being weighted with weight  $w_j$
- The total sum is  $h$
- If  $h$  is more than some threshold  $\theta$ 
  - then neuron fires:  $y = 1$ ,
  - else not:  $y = 0$  (sometimes also used:  $y = -1$ )

# Perceptron Neurons

$n$  input neurons

$$h = \sum_{j=1}^n x_j w_j$$

$$y = \begin{cases} 1 & h \geq \theta \\ 0 & h < \theta \end{cases}$$

for some threshold  $\theta$

*How biologically (un)realistic?*


- The weight  $w_j$  can be positive or negative
- A unit can become inhibitory or excitatory, or both
- Use only a linear sum of inputs
- Use a simple output instead of a pulse (spike train)
- No refractory period

# Some Terminology

<i>Term</i>	<i>Typical Symbol</i>	<i>Alternate Term(s)</i>
■ Input vector	$\mathbf{x}$	<i>input activation</i>
■ Weights	$w_{ij}$	<i>synaptic weights (from <math>j</math> to <math>i</math>)</i>
■ Inner activation	$h$	<i>net activation</i>
■ Activation function	$g$	<i>transfer function; threshold function</i>
■ Output	$y$	<i>(outer) activation; prediction</i>
■ Target	$t$	<i>teacher value</i>
■ Error	$E$	<i>cost</i>




# Weight Space: Represent a Unit with its Incoming Weights



# Neural Networks

- Started by psychologists and neurobiologists as computational analogues of neurons
- A neural network: A set of connected input/output units where each connection has a **weight** associated with it
- During supervised learning, the **network learns by adjusting the weights** so as to be able to predict the correct class label of the input tuples
- Also referred to as **connectionist learning** due to the connections between units

# Perceptron Network



# Updating the Weights

$$w_{ij} \leftarrow w_{ij} + \Delta w_{ij}$$

- We want to change the values of the weights
- Aim: **minimize the error** at the output
- Let error  $E = t - y$ . We want  $E$  to be 0
- Use:

Diagram illustrating the weight update formula:

$$\Delta w_{ij} = \eta \cdot (t_i - y_i) \cdot x_j$$

The components of the formula are labeled with red arrows:

- $\Delta w_{ij}$ : Change of weight
- $\eta$ : Learning rate
- $(t_i - y_i)$ : Error
- $x_j$ : Input

# Perceptron Algorithm

- Initialisation: set all weights to small positive and negative random numbers
- For #iterations
  - Chose a new data point  $(\mathbf{x}, t)$
  - Compute the output activation  $y_i$  of each neuron  $i$


$$h_i = \sum_{j=1}^n w_{ij} x_j \quad y_i = \begin{cases} 1 & h_i \geq \theta \\ 0 & h_i < \theta \end{cases}$$

- Update each of the weights according to



$$\Delta w_{ij} = \eta \cdot (t_i - y_i) \cdot x_j$$

$$w_{ij} \leftarrow w_{ij} + \Delta w_{ij}$$

# Perceptron Network



# Obstacle Avoidance with the Perceptron



$$\eta = 0.3$$
$$\theta = -0.01$$
$$w_{ij} = 0$$


# Obstacle Avoidance with the Perceptron:

## Behaviour we want

	LS	RS	LM	RM	
	<b>0</b>	<b>0</b>	<b>1</b>	<b>1</b>	
	0	1	-1	1	
	1	0	1	-1	
	1	1	X	X	



# Obstacle Avoidance with the Perceptron



Assume initial weights are 0  
No update if target = actual computed

$$\Delta w_{ij} = \eta \cdot (t_i - y_i) \cdot x_j$$

$$w_{ij} \leftarrow w_{ij} + \Delta w_{ij}$$

$$w_{11} = 0 + 0.3 \cdot (1 - 1) \cdot 0 = 0$$


$$w_{21} = 0 + 0.3 \cdot (1 - 1) \cdot 0 = 0$$

And the same for  $w_{12}$ ,  $w_{22}$

# Obstacle Avoidance with the Perceptron

	<b>LS</b>	<b>RS</b>	<b>LM</b>	<b>RM</b>
	0	0	1	1
	<b>0</b>	<b>1</b>	<b>-1</b>	<b>1</b>
	1	0	1	-1
	1	1	X	X

# Obstacle Avoidance with the Perceptron



$w_{12}$ : the robot turns left by reversing the left motor

No update if input = 0

$$\Delta w_{ij} = \eta \cdot (t_i - y_i) \cdot x_j$$

$$w_{ij} \leftarrow w_{ij} + \Delta w_{ij}$$

$$w_{11} = 0 + 0.3 \cdot (-1 - 1) \cdot 0 = 0$$

$$w_{21} = 0 + 0.3 \cdot (1 - 1) \cdot 0 = 0$$


$$w_{12} = 0 + 0.3 \cdot (-1 - 1) \cdot 1 = -0.6$$

$$w_{22} = 0 + 0.3 \cdot (1 - 1) \cdot 1 = 0$$

# Obstacle Avoidance with the Perceptron

LS	RS	LM	RM
0	0	1	1
0	1	-1	1
1	0	1	-1
1	1	X	X

# Obstacle Avoidance with the Perceptron



$$\Delta w_{ij} = \eta \cdot (t_i - y_i) \cdot x_j$$

$$w_{ij} \leftarrow w_{ij} + \Delta w_{ij}$$


$$w_{11} = 0 + 0.3 \cdot (1 - 1) \cdot 1 = 0$$

$$w_{21} = 0 + 0.3 \cdot (-1 - 1) \cdot 1 = -0.6$$


$$w_{12} = -0.6 + 0.3 \cdot (1 - 1) \cdot 0 = -0.6$$

$$w_{22} = 0 + 0.3 \cdot (-1 - 1) \cdot 0 = 0$$

# Obstacle Avoidance with the Perceptron



# Obstacle Avoidance with a Mindstorm Vehicle



# Linear Separability

■ Outputs are:  $y_i = \text{sign}\left(\sum_{j=1}^n w_{ij} x_j\right)$

- Positive output +1 if:

$$w_i \cdot x \geq 0$$


- Negative output -1 (or 0) if:

$$w_i \cdot x < 0$$

weight  
vector

dot  
product

input  
vector




The region in input space where  $x$  yield positive output  $y$  is a half-plane.




# Bias

- An extra input – increases or lowers the net input (depending on its sign)
- Can be regarded as a weight connected to a **constant of 1**
  - Then bias learning is similar to weight learning
  - Can **convert** a threshold into an additional weight.




# Perceptron with Bias




- The  $n$ -dimensional input vector  $\mathbf{x}$  together with bias  $b$  is mapped into variable  $y$  by means of the scalar product and a nonlinear function mapping
- negative *bias*  $\sim$  (positive) *threshold*

# Linear Separability




# Linear Separability



The Binary  
AND Function

# Limitations of the Perceptron


Linear Separability?




Exclusive Or (XOR) function

A	B	Out
0	0	0
0	1	1
1	0	1
1	1	0

# Limitations of the Perceptron




# Limitations of the Perceptron



One way around the problem is to use a more complex input set (e.g., three-dimensional:  $x_3 = x_1 \cdot x_2$  ).

Another is to make the network more complex.

# Perceptrons – Early Successes(?)







# Perceptron

- How can we make the perceptron more powerful?
- More layers in the networks?
- More connections?
- Perceptron: one layer of weights
- Multi-layer perceptron: at least 2 layers of weights

# The Multi-Layer Perceptron



# XOR Again




So E does not fire

C does not fire,  
D does not fire

Apply input 0 0

# XOR Again




So overall E fires

C fires  
D does not fire

Apply input 1 0

Same for input 0 1


# XOR Again





So E does not fire

C fires  
D fires

Apply input 1 1




XOR Again



# Gradient Descent

- The MLP *can* solve XOR
- How do we choose the weights?
- Harder than for the perceptron
  - More weights
  - Which weights are wrong? Input-hidden or hidden-output?
- Use gradient descent learning
- Compute gradient  $\Rightarrow$  differentiation

# Gradient Descent



If we differentiate function E, we get the negative gradient of the function (direction of change)



# An Error Function

- If  $E=(t-y) \rightarrow$  pos. and neg. errors would cancel out
- Better: **sum-of-squares error**

$$E(\mathbf{w}) = \frac{1}{2} \sum_i (t_i - y_i)^2 = \frac{1}{2} \sum_i \left( t_i - \sum_j w_{ij} x_j \right)^2$$

- We will ignore the threshold function in the output neurons

$$\Rightarrow -\frac{\partial E}{\partial w_{ij}} = (t_i - y_i) \cdot x_j \quad \text{Gradient descent}$$

- Rule for the weights to the output layer (also for perceptron)

# A Multi-Layer Feed-Forward Neural Network


Output vector

Output layer

Hidden layer

Input layer

Input vector:  $X$



# How a Multi-Layer Neural Network Works

- The **inputs** to the network correspond to the attributes measured for each training tuple
- Inputs are fed simultaneously into the units of the **input layer**
- They are then weighted and fed simultaneously to a **hidden layer**
- The weighted outputs of the last hidden layer are input to units making up the **output layer**, which emits the network's prediction
- The network is **feed-forward**, i.e. none of the weights cycles back to a unit in the same or a previous layer
- From a statistical point of view, networks perform **nonlinear regression**: given enough hidden units and enough training samples, they can closely approximate any function

# Decide on the Network Topology

- # of units in the *input layer* ←fixed through the application
  - One input unit per domain value
- # of *hidden layers* (if  $> 1$ )
  - Complex function – transformations? Hierarchical features?
- # of units in *each hidden layer*
  - Complex function – many features?
- # of units in the *output layer* ←fixed through the application
  - One output unit for each variable in regression
  - A single output unit for two-class classification
  - For more than two classes, one output unit per class
    - output values may be coupled by a softmax function
- *May repeat training with different network topologies*


# Weight Initialisation & Data Preprocessing

- Initialize the network with small random weights
  - All-same weights would lead to symmetry-breaking problem: all hidden units will have same activations and learn the same
  - Large weights could lead to saturation of the transfer function
- Normalizing the input values for each attribute measured in the training tuples, e.g.
  - shift & scale attribute values to be in the interval  $[0.0 \dots 1.0]$ , or
  - shift & scale them to have mean=0, variance=1;
  - this may be done **per attribute** or **over all attributes**
    - all attributes have same importance*
    - attributes keep their relative importance*
- *May repeat training with a different set of initial weights*

# Training MLP

## (1) Forward Pass

- Put the input values in the input layer
- Calculate the activations of the hidden nodes
- Calculate the activations of the output nodes
- Calculate the errors using the targets




For example

$$t_i - y_i$$

# Training MLPs

## (2) Backward Pass

- Using output errors, update last layer of weights
- Calculate hidden-layer errors, update hidden-layer weights
- Work backwards through the network
- Error is backpropagated through the network



# Error Backpropagation

- Iteratively process training tuples & compare the network's prediction with the actual known target value
- For each training tuple, the weights are modified to **minimize** the **squared error** between network's prediction and actual target value
  - This minimizes the **mean** square error over the entire data set
- Errors are computed “**backwards**”: from the output layer, through each hidden layer down to the first hidden layer, hence “**backpropagation**”
- Steps
  - Initialize weights (to small random #s) and biases in the network
  - For each data point:
    - Propagate the inputs forward (by applying activation function)
    - Propagate the error backwards (backpropagation)
    - Update the weights and biases (using inputs and errors)
  - Terminating condition (when error small; test error increases; etc.)



# Activation Function


- In the analysis we ignored the activation function
  - The threshold function is not differentiable
- What do we want in an activation function?
  - Differentiable
  - Should saturate (become constant at ends)
  - Change between saturation values quickly

# Sigmoid Neurons

- Sigmoidal / logistic transfer function:
  - gives a real-valued, positive output
  - bounded in interval  $[0,1]$
  - easily differentiable, positive derivative
  - output can be interpreted as a **probability** of a binary output to be =1 (or of producing a spike) → stochastic binary neurons

$$h = b + \sum_j x_j w_j$$

$$y = g(h) = \frac{1}{1 + e^{-h}}$$




# Sigmoid Activation Function for a Neuron

Transfer function:

$$g(h) = \frac{1}{1 + \exp(-h)}$$

Derivative:

$$\begin{aligned} g'(h) &= \frac{\partial g(h)}{\partial h} \\ &= \dots \\ &= g(h) \cdot (1 - g(h)) \end{aligned}$$



The derivative can be expressed as a function of the **outputs**.

# Overview of Transfer Functions

## Transfer function:

- sigmoid

$$g(h) = \frac{1}{1 + \exp(-h)}$$

- linear

$$g(h) = h$$

- threshold function

$$g(h) = \begin{cases} 1 & h \geq \theta \\ 0 & h < \theta \end{cases}$$

- sign

$$g(h) = \begin{cases} 1 & h \geq \theta \\ -1 & h < \theta \end{cases}$$

## Corresponding derivative:

$$g'(h) = g(h) \cdot (1 - g(h))$$

$$g'(h) = 1$$

no useful derivative

no useful derivative

# Error Terms

- Need to differentiate the sigmoid function
- Gives us the following **error terms** (deltas)
  - For the outputs

$$\delta_i = (t_i - y_i) \underbrace{y_i(1 - y_i)}_{\text{derivative}}$$

- For the hidden nodes (with activations  $y_j^{hid}$ )

$$\delta_j = \underbrace{y_j^{hid}(1 - y_j^{hid})}_{\text{derivative}} \sum_i w_{ij} \delta_i$$

# Update Rules

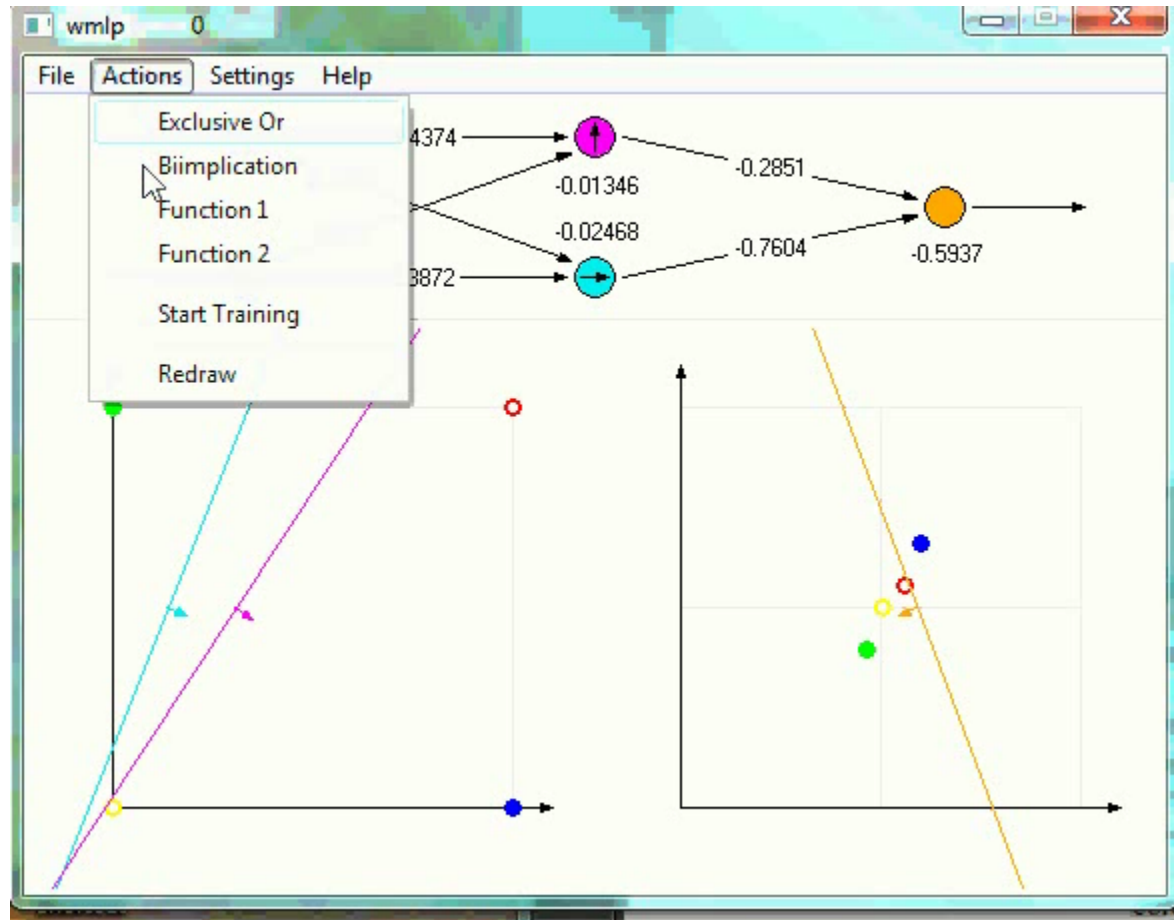
- This gives us the necessary update rules
  - For the weights connected to the outputs:

$$w_{ij} \leftarrow w_{ij} + \eta \delta_i y_j^{\text{hid}}$$

- For the weights connected to the hidden nodes:

$$v_{jk} \leftarrow v_{jk} + \eta \delta_j x_k$$

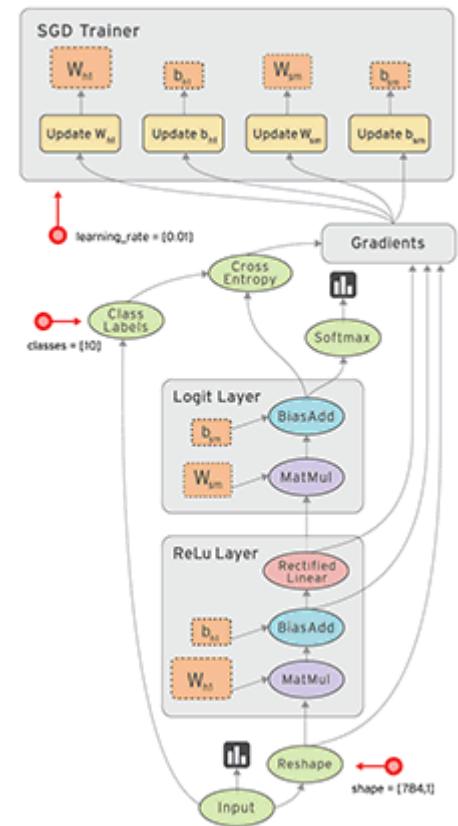
# MLP training a XOR problem



[<http://www.borgelt.net/mlpd.html>]

# Tensorflow

- Open source package for deep MLP learning by google
- Given a network structure and cost function:
  - does automatic differentiation and learning
- Online demo for small networks:
  - <http://playground.tensorflow.org>



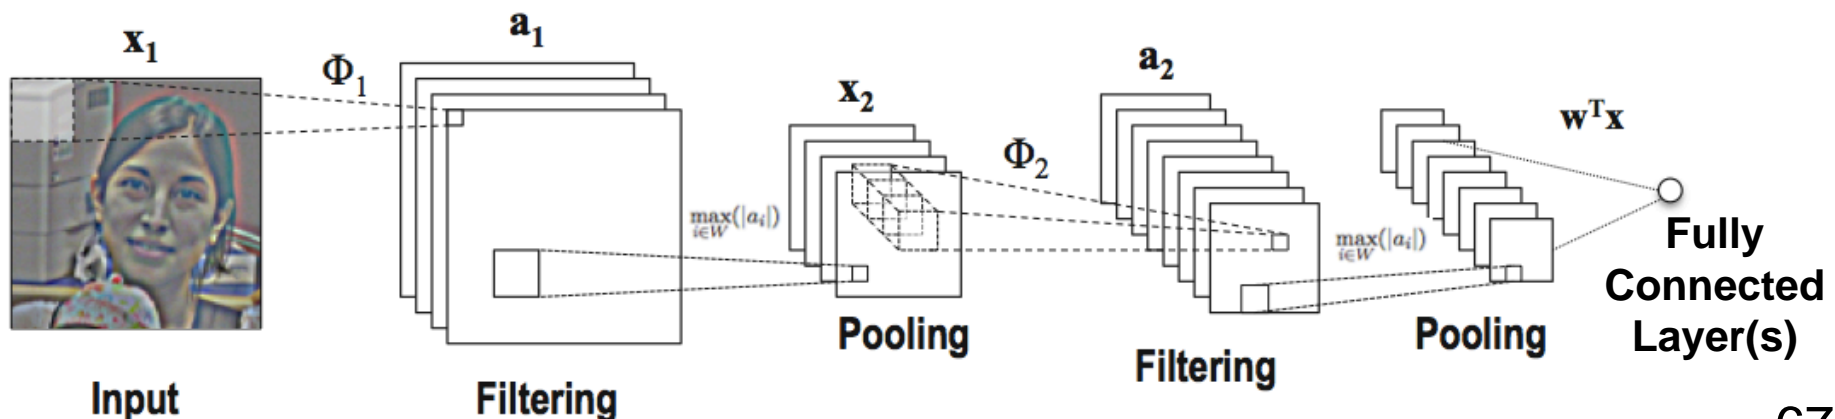


# Network Topology

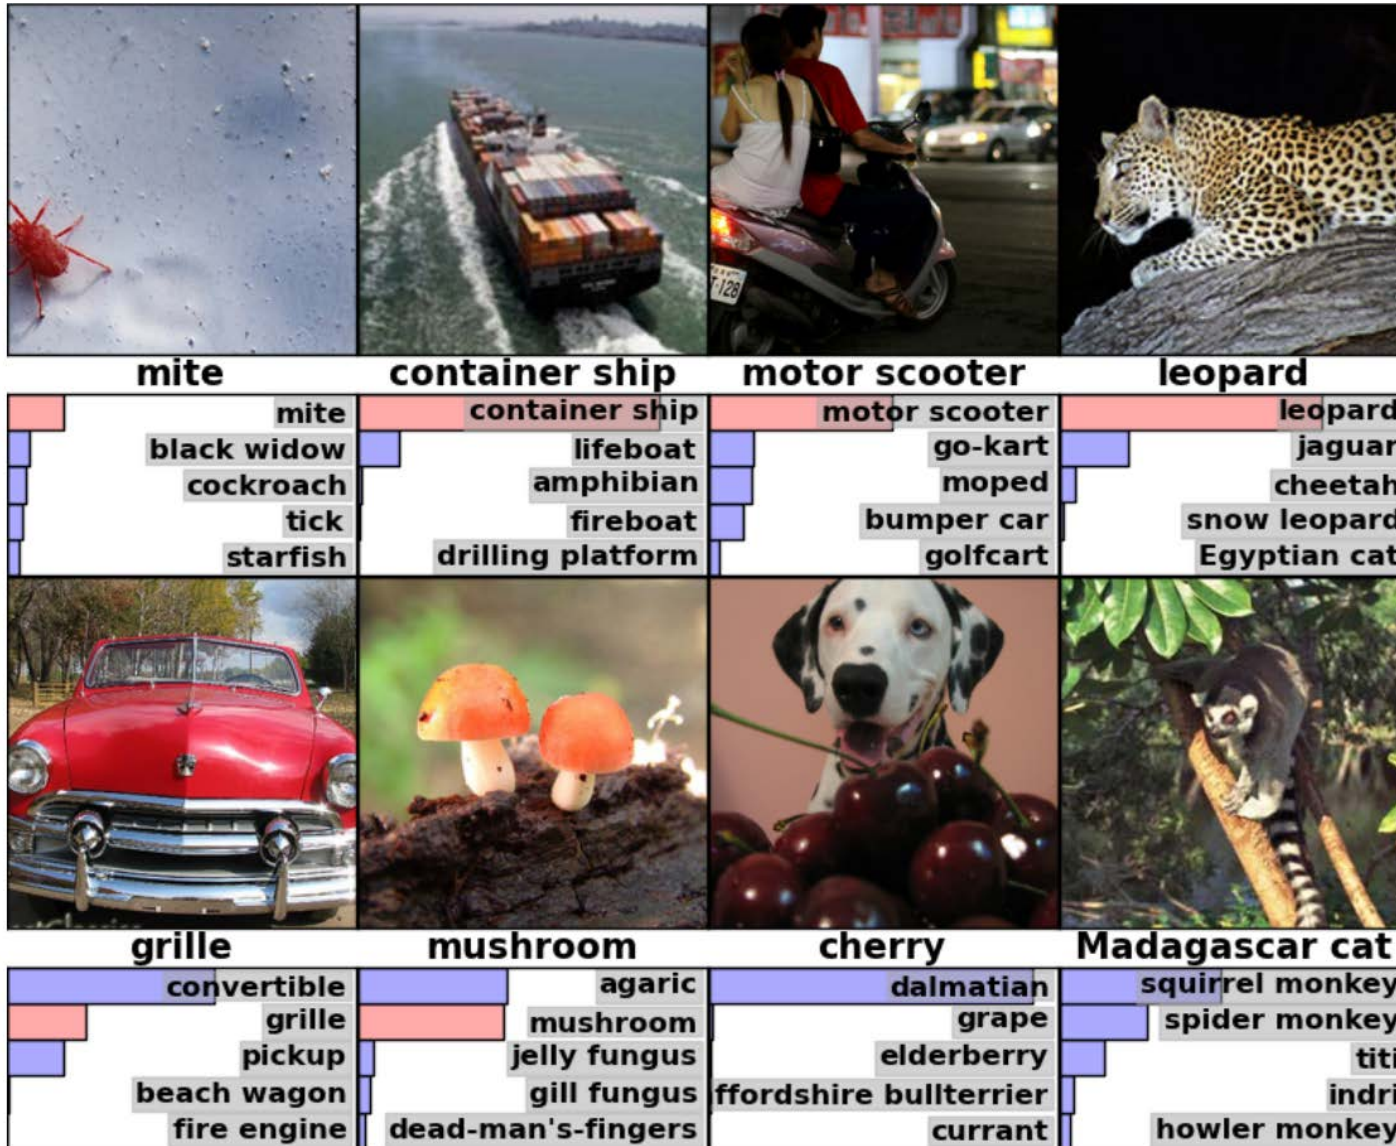
- How many layers?
- How many neurons per layer?
- Experiments
  - Often two or three hidden layers (but new research into deep learning networks...)
  - Determine size of layers (usually get smaller)
  - Test several different networks

# Deep Learning (MLP with many Layers)

- Enabled by **GPUs**, multi-core CPUs and **large data sets**
- Convolutional **filtering** layers (only lower layers)
  - Use small, replicated (shared) weights on large inputs
  - Lowest-layer filters sometimes not learned  
→ e.g. Sobel-, Gabor-, centre-surround/whitening filters
- MAX-**Pooling** layers (over convolutional layers)
  - Create invariances (e.g. to shift/scale); reduce dimensionality




# Deep Learning Multi-Class Classification



# Batch and Incremental Learning

- When should the weights be updated?
  - After all inputs seen (*batch*)
    - Accurate estimate of gradient
    - Converges systematically to the (local) minimum
    - Requires many epochs (passes through the whole dataset)
  - After each input is seen (*incremental, online*)
    - Simpler to program
    - Handles infinite amount of data (continual learning)
    - Noise may help escaping from saddle points in the energy landscape, or even from local minima
    - Pitfall: data distribution may drift (also within batch data).  
Remedy: randomize order of presentation

# Gradient Descent




- Local gradient does not point towards minimum
- Gradient descent with large learning rate  $\rightarrow$  oscillations
- Long learning time!

# Gradient Descent


- Learning rule (for the simple case of the perceptron):

$$-\frac{\partial E}{\partial w_{ij}} = (t_i - y_i) \cdot x_j$$



- Assume: inputs  $x_1$  and  $x_2$  are of similar importance for classification
  - both have mean zero:  $\mu(x_1) = \mu(x_2) = 0$
  - but std. deviations differ:  $\sigma(x_1) < \sigma(x_2)$
- weights should be:  $w_1 > w_2$   
but average updates:  $|\Delta w_1| < |\Delta w_2|$

# Momentum Term




$$w_{ij}^{\tau} \leftarrow w_{ij}^{\tau-1} + \eta \delta_i y_j^{\text{hid}} + \underbrace{\alpha \Delta w_{ij}^{\tau-1}}$$

Add contribution from previous weight change (**momentum**)


- Counteracts oscillations, by averaging previous and current updates (relevant for batch learning)
- Averages out noise (relevant for on-line learning)
- More stable, leads to faster learning



# Learning Capacity




Output of one sigmoid



Output of two sigmoids




# Learning Capacity




Addition of two ridges  
Unique maximum

Addition of more ridges  
and transformation with  
another sigmoid  
Localised response




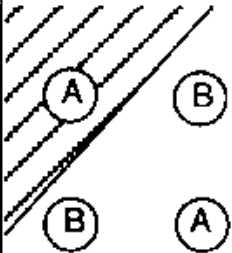
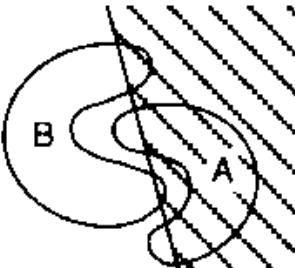


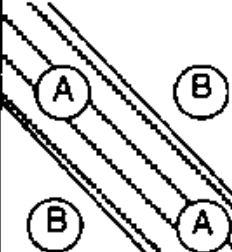
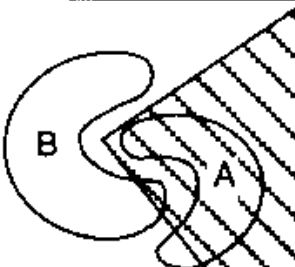


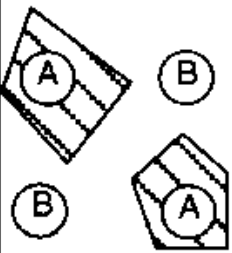
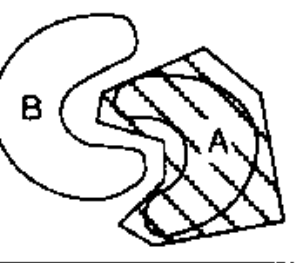
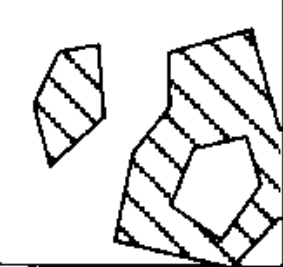
# Learning Capacity

Inputs



**Any function can be approximated as the summation of many localised responses**

# Decision Boundaries (Lippmann)

Structure	Types of Decision Regions	Exclusive OR Problem	Classes with Meshed Regions	Most General Region Shapes
Single-Layer 	Half Plane Bounded by Hyperplane			
Two-Layer 	Convex Open or Closed Regions			
Three-Layer 	Arbitrary (Complexity Limited by Number of Nodes)			



# Generalisation

- Aim of neural network learning:

Generalise from training examples to all possible inputs

- Undertraining is bad
- Overtraining is worse

# Overfitting

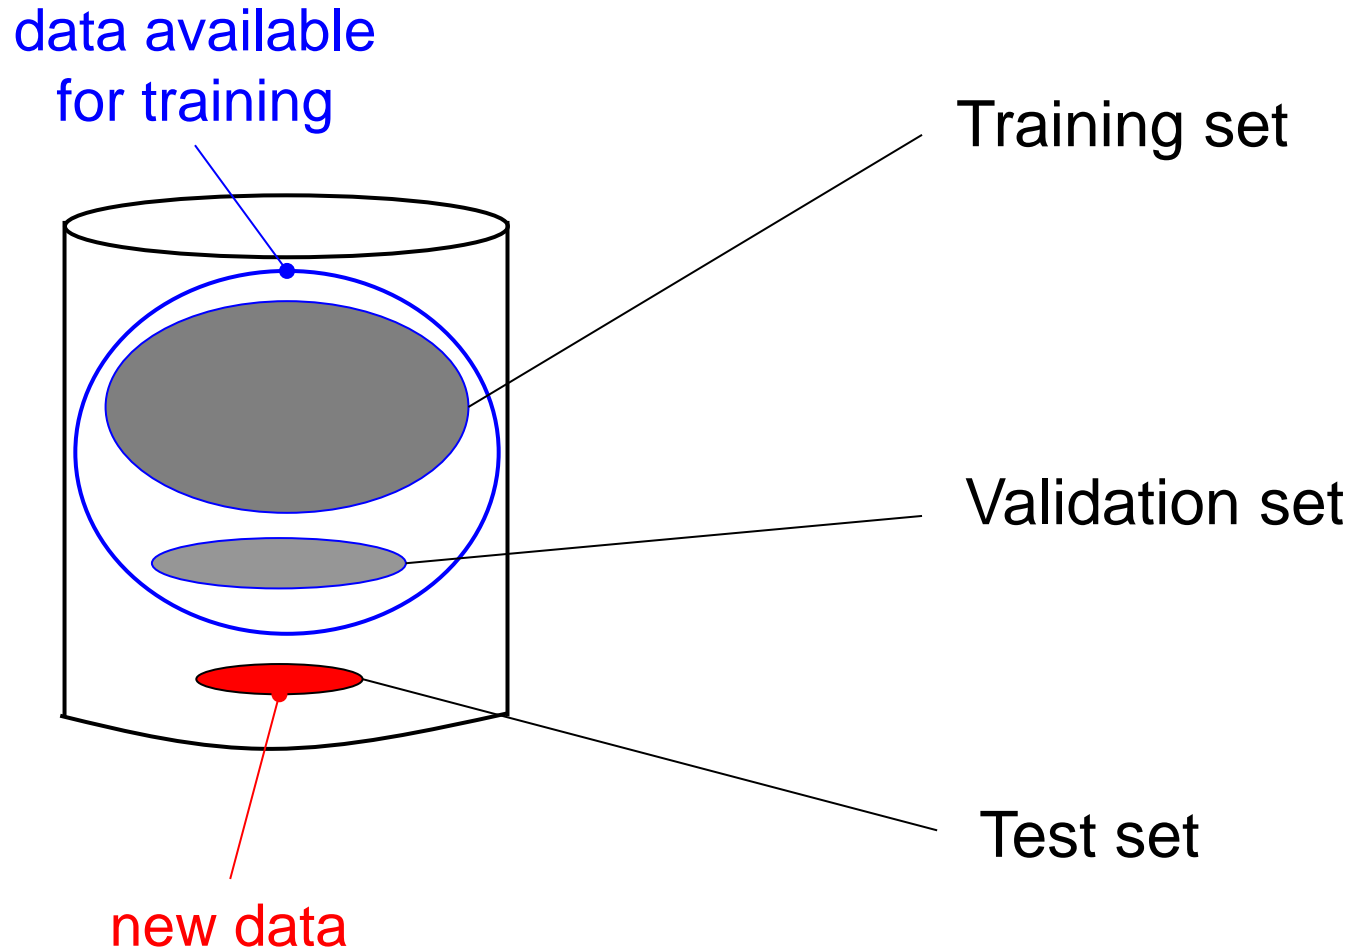


# Testing

How do we evaluate our trained network?

- The error on the **training data** is biased and hides overfitting
- Validate on a separate **validation set**
  - evaluate periodically on this validation set during training (while training only on the training set)
  - indicator of overfitting: the validation error increases
- After training, test the final model on the **test set**

# Using Training, Validation and Test Data

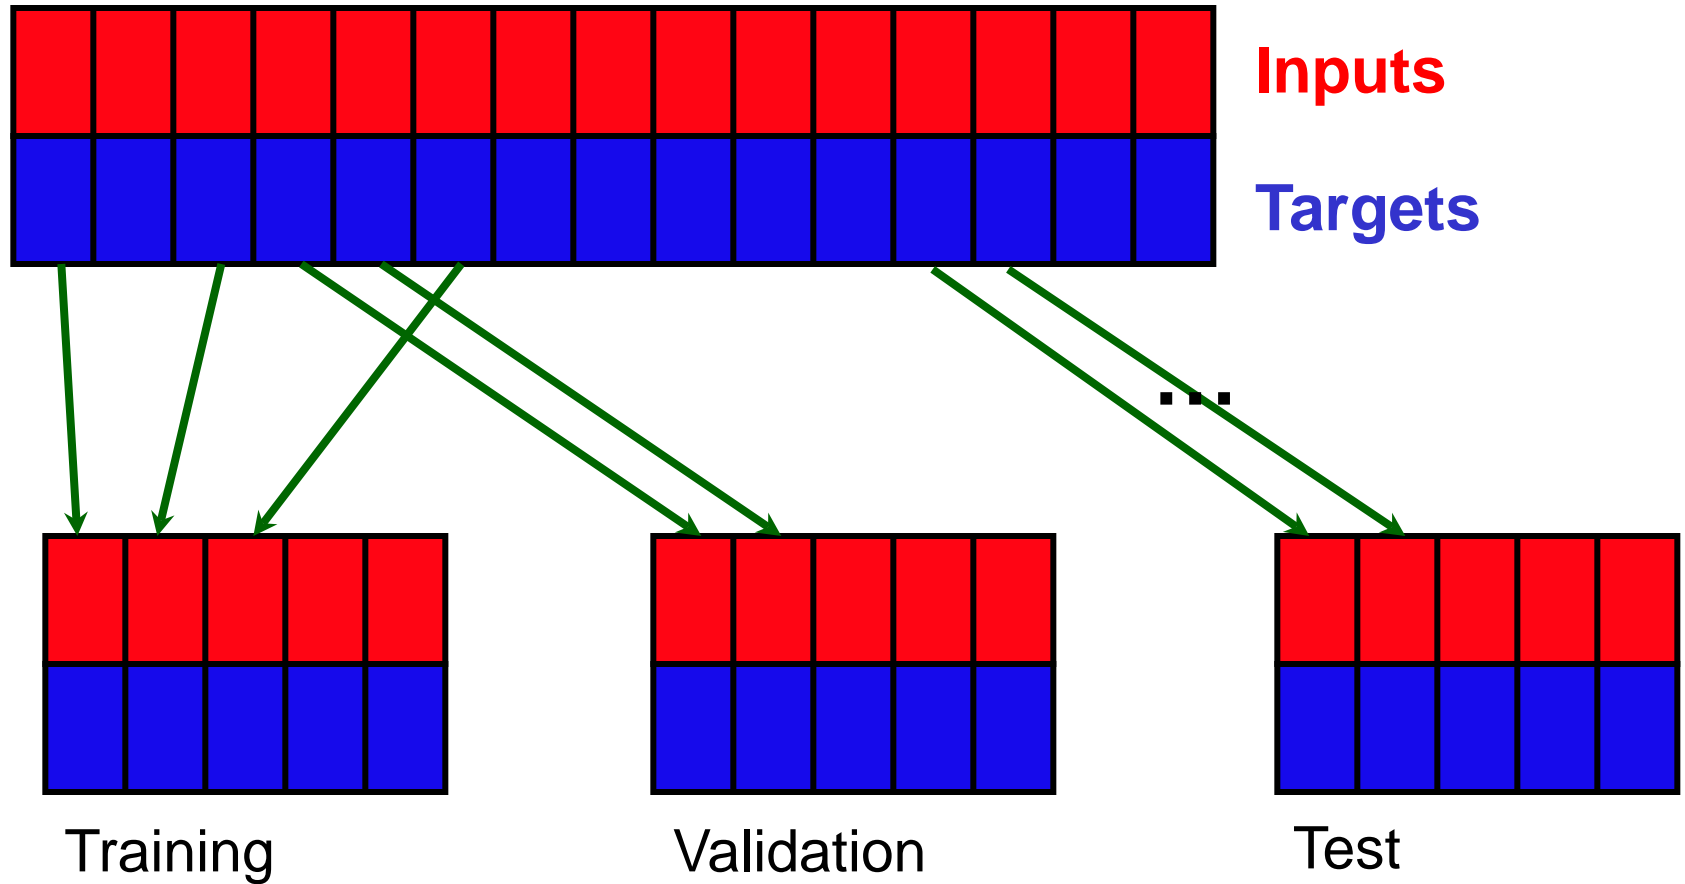


# Validation

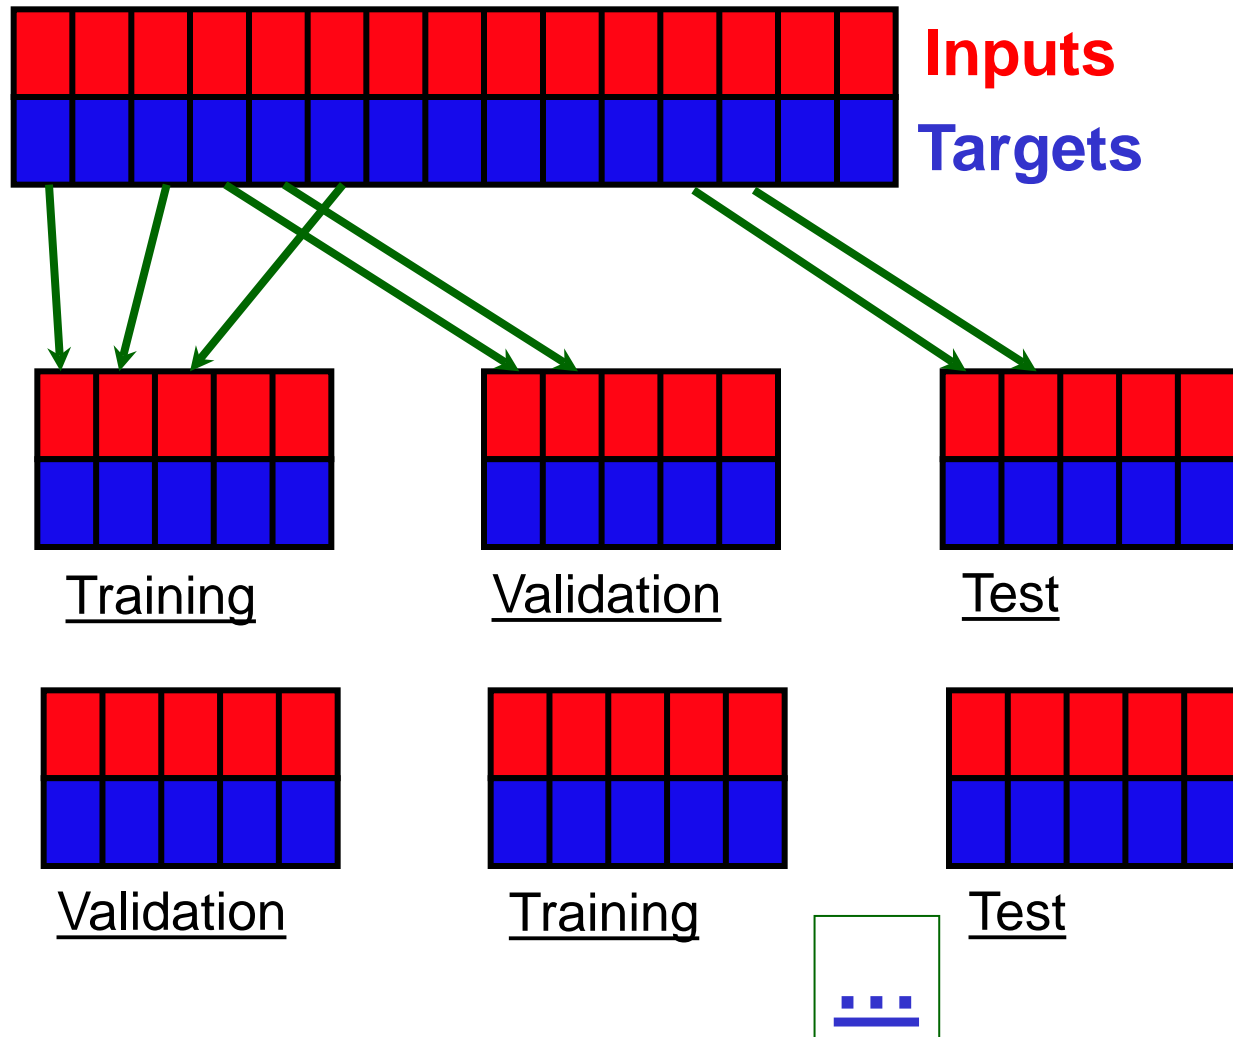
- Of the data that is available during training, keep a subset for validation
- Train the network on **training** data
- Periodically, stop and evaluate on **validation** set
- After training has finished, test on **test** set
- This is coming expensive on data!



# Hold Out Cross Validation



# Multifold Cross Validation



# Evaluating Classifier Accuracy: Holdout & Cross-Validation Methods

## ■ **Holdout method**

- Given data is randomly partitioned into two independent sets
  - Training set (e.g., 2/3) for model construction
  - Test (“hold-out”) set (e.g., 1/3) for accuracy estimation
- Expensive on the data

## ■ **Cross-validation** ( $k$ -fold, where $k = 10$ is popular)


- Randomly partition the data into  $k$  *mutually exclusive* subsets, each approximately equal size
- At  $i$ -th iteration, use  $D_i$  as test set and others as training set
- Leave-one-out:  $k$  folds where  $k = \#$  of tuples, for small sized data
- Random subsampling:  $k$  folds, with *random split* between training and test set each time; accuracy = avg. of the accuracies obtained

# Early Stopping

When should we stop training?

- Could set a minimum training error
  - Danger of overfitting
- Could set a number of epochs
  - Danger of underfitting or overfitting
- Can use the validation set
  - Measure the error on the validation set during training

# Early Stopping



# Backpropagation and Interpretability

- **Rule extraction** from networks: network pruning
  - Simplify the network structure by removing weighted links that have the least effect on the trained network
  - Then perform link, unit, or activation value clustering
  - The set of input and activation values are studied to derive rules describing the relationship between the input and hidden unit layers
- **Sensitivity analysis**: assess the impact that a given input variable has on a network output.
  - The knowledge gained from this analysis can be represented in rules

# Summary: Neural Networks as a Classifier

## ■ Weaknesses

- Long training time (but same with humans ...)
- Some parameters have to be determined empirically, e.g.
  - network topology, transfer functions, learning rate, etc.
- Challenging to interpret the symbolic meaning behind the learned weights and of “hidden units”
- Cannot handle well missing values

## ■ Strengths

- Well-suited for continuous-valued inputs and outputs
- High tolerance to noisy data
- Generalisation ability: classify untrained patterns
- Successful on a wide array of real-world data
- Algorithms are inherently parallel
- Recent techniques extract rules from trained neural networks
- Relationship to brain

# WTM Student Project



(a) Learner and teacher



(b) Objects

In: IEEE First International Conference on Cognitive Systems and Information Processing (CSIP 2012), in press (Springer). Beijing, December 15-17, 2012.


## Object Learning with Natural Language in a Distributed Intelligent System – A Case Study of Human-Robot Interaction

Stefan Heinrich, Pascal Folleher, Peer Springstübe, Erik Strahl,  
Johannes Twiefel, Cornelius Weber, and Stefan Wermter

University of Hamburg, Department of Informatics, Knowledge Technology  
Vogt-Kölln-Straße 30, D - 22527 Hamburg, Germany  
{heinrich,6follehe,3springs,strahl}@informatik.uni-hamburg.de  
{7twiefel,weber,wermter}@informatik.uni-hamburg.de  
<http://www.informatik.uni-hamburg.de/WTM/>



# Student Project: Classifying Objects with MLPs



The robot perceives visual features of the objects and learns the objects' names

[http://www.informatik.uni-hamburg.de/WTM/teaching/WiSe11\\_HumanRobotInteraction\\_Pj.shtml](http://www.informatik.uni-hamburg.de/WTM/teaching/WiSe11_HumanRobotInteraction_Pj.shtml)