

Machine Learning: A General View

Dr. Víctor Uc Cetina

Facultad de Matemáticas
Universidad Autónoma de Yucatán

`cetina@informatik.uni-hamburg.de`
`https://sites.google.com/view/victorucetina`

Content

- 1 Machine Learning
- 2 About this ML course
- 3 Supervised Learning
- 4 Unsupervised Learning
- 5 Reinforcement Learning

What is Machine Learning?

- Machine learning (ML) refers to algorithms used to extract patterns from data and learn a mathematical model that could be used by a computer program to make intelligent decisions.

Some Data

```
010101010100  
010010101010  
101000101010  
100101001100  
100101000111  
0...
```

A Model



$$y = f(x)$$

Decision Making

```
if y=a then do A  
else if y=b then do B  
else if y=c then do C  
...
```

What is Machine Learning?

- Example 1. Given a set of digital images of blood samples containing Chagas parasites, decide if a digital image of a new blood sample contains at least a Chagas parasite.



What is Machine Learning?

- Example 2. Given a set of text phrases written in two different languages, decide if a new text phrase belongs to one of the languages existing in your set of phrases.

hola a todo el mundo - no me gusta decir
adiós - hello people - adoro la comida - our
world is wonderful - el planeta agua -
ciencia ficción es ciencia - el algoritmo más
rápido - la mesa es redonda - the door is
black -
my chair is broken - el plato está limpio,
esa escalera está muy inclinada - my
mouse is wireless - an electronic book - a
wide road is better, we were at home - ...

What is Machine Learning?


- Example 3. Given a history of the commands used to control a drone, decide which is the best command to perform in order to avoid a collision with the ground.




What is Machine Learning?

- Using machine learning algorithms usually involves at least two phases:

LEARNING PHASE



APPLICATION PHASE



What is Machine Learning?

- Based on the type of data available and the decisions needed, we can talk about three general kinds of machine learning algorithms:
 - Supervised Learning (SL)
 - Unsupervised Learning (UL)
 - Reinforcement Learning (RL)

What is Machine Learning?

- People involved in ML research and development comes from fields such as: Computer Science, Artificial Intelligence, Statistics, Neuroscience, Psychology, Robotics, Bioinformatics, etc.

About Learning Machine Learning

Main Human Learning Loop for ML:

- 1 Understanding the kind of machine learning problem (SL, UL, RL)

About Learning Machine Learning

Main Human Learning Loop for ML:

- 1 Understanding the kind of machine learning problem (SL, UL, RL)
- 2 Understanding the mathematics (basic statistics, probability, calculus, linear algebra)

About Learning Machine Learning

Main Human Learning Loop for ML:

- 1 Understanding the kind of machine learning problem (SL, UL, RL)
- 2 Understanding the mathematics (basic statistics, probability, calculus, linear algebra)
- 3 Understanding the algorithms (data structures, complexity)

About Learning Machine Learning

Main Human Learning Loop for ML:

- 1 Understanding the kind of machine learning problem (SL, UL, RL)
- 2 Understanding the mathematics (basic statistics, probability, calculus, linear algebra)
- 3 Understanding the algorithms (data structures, complexity)
- 4 Coding (Matlab/Octave)

About Learning Machine Learning

Main Human Learning Loop for ML:

- 1 Understanding the kind of machine learning problem (SL, UL, RL)
- 2 Understanding the mathematics (basic statistics, probability, calculus, linear algebra)
- 3 Understanding the algorithms (data structures, complexity)
- 4 Coding (Matlab/Octave)
- 5 Experimenting (running programs and plotting graphs)

Important to know

- Course materials will be available to download from the Stine website

Important to know

- Course materials will be available to download from the Stine website
- Feedback of exercises are given in person during the exercise sessions or in my office.

Important to know

- Course materials will be available to download from the Stine website
- Feedback of exercises are given in person during the exercise sessions or in my office.
- You need to complete and deliver at least 5 exercises to have the right to write the final exam.

Part I - Supervised Learning

① Linear regression

Part I - Supervised Learning

- 1 Linear regression
- 2 Classification

Part I - Supervised Learning

- ① Linear regression
- ② Classification
- ③ Generative learning models

Part I - Supervised Learning

- 1 Linear regression
- 2 Classification
- 3 Generative learning models
- 4 Support vector machines

Part I - Supervised Learning

- 1 Linear regression
- 2 Classification
- 3 Generative learning models
- 4 Support vector machines
- 5 AdaBoost

Part I - Supervised Learning

- 1 Linear regression
- 2 Classification
- 3 Generative learning models
- 4 Support vector machines
- 5 AdaBoost
- 6 Learning theory

Part I - Supervised Learning

- 1 Linear regression
- 2 Classification
- 3 Generative learning models
- 4 Support vector machines
- 5 AdaBoost
- 6 Learning theory
- 7 Regularization and model selection

Part II - Unsupervised Learning

1 Expectation-Maximization (EM) algorithm

Part II - Unsupervised Learning

- 1 Expectation-Maximization (EM) algorithm
- 2 Factor analysis

Part II - Unsupervised Learning

- 1 Expectation-Maximization (EM) algorithm
- 2 Factor analysis
- 3 Principal component analysis

Part III - Reinforcement Learning

1 Dynamic programming

Part III - Reinforcement Learning

- 1 Dynamic programming
- 2 Montecarlo methods

Part III - Reinforcement Learning

- 1 Dynamic programming
- 2 Montecarlo methods
- 3 Temporal difference learning

Part III - Reinforcement Learning

- 1 Dynamic programming
- 2 Montecarlo methods
- 3 Temporal difference learning
- 4 Generalization and function approximation

Sources of Information

- Presentations

Sources of Information

- Presentations
- ML notes written by Prof. Andrew Ng

Sources of Information

- Presentations
- ML notes written by Prof. Andrew Ng
- **Pattern Recognition and Machine Learning.** Christopher M. Bishop. Springer. 2006.

Sources of Information

- Presentations
- ML notes written by Prof. Andrew Ng
- **Pattern Recognition and Machine Learning.** Christopher M. Bishop. Springer. 2006.
- **Foundations of Machine Learning.** Mehryar Mohri, Afshin Rostamizadeh and Ameet Talwalkar. The MIT Press. 2012.

Sources of Information

- Presentations
- ML notes written by Prof. Andrew Ng
- **Pattern Recognition and Machine Learning.** Christopher M. Bishop. Springer. 2006.
- **Foundations of Machine Learning.** Mehryar Mohri, Afshin Rostamizadeh and Ameet Talwalkar. The MIT Press. 2012.
- **Understanding Machine Learning: from Theory to Algorithms.** Shai Shalev-Shwartz and Shai Ben-David. Cambridge University Press. 2014.

Sources of Information

- Presentations
- ML notes written by Prof. Andrew Ng
- **Pattern Recognition and Machine Learning**. Christopher M. Bishop. Springer. 2006.
- **Foundations of Machine Learning**. Mehryar Mohri, Afshin Rostamizadeh and Ameet Talwalkar. The MIT Press. 2012.
- **Understanding Machine Learning: from Theory to Algorithms**. Shai Shalev-Shwartz and Shai Ben-David. Cambridge University Press. 2014.
- **Reinforcement Learning: An Introduction (second edition)**. Richard Sutton and Andrew Barto. The MIT Press. 2018.

Regression

- In the **regression** problem

we have

$$R : X \rightarrow Y,$$

$$X \subseteq \mathbb{R}^d,$$

$$Y \subseteq \mathbb{R}. \text{ (Continuous)}$$

Example:

$$y =$$

$$w_0 + w_1x + w_2x^2 + w_3x^3$$

Regression

- In the **regression** problem

we have

$$R : X \rightarrow Y,$$


$$X \subseteq \mathbb{R}^d,$$

$$Y \subseteq \mathbb{R}. \text{ (Continuous)}$$

Example:

$$y =$$

$$w_0 + w_1x + w_2x^2 + w_3x^3$$



Regression

- In the **regression** problem

we have

$$R : X \rightarrow Y,$$


$$X \subseteq \mathbb{R}^d,$$

$$Y \subseteq \mathbb{R}. \text{ (Continuous)}$$

Example:

$$y =$$

$$w_0 + w_1x + w_2x^2 + w_3x^3$$



Classification

- In the **classification**

problem we have

$$C : X \rightarrow Y,$$

$$X \subseteq \mathbb{R}^d,$$

$$Y = \{y_1, y_2, \dots, y_n\}.$$

(Discrete)

Example:

$$y = \frac{1}{1 + e^{-(w_0 + w_1 x_1 + w_2 x_2)}}$$

Classification

- In the **classification** problem we have

$$C : X \rightarrow Y,$$


$$X \subseteq \mathbb{R}^d,$$

$$Y = \{y_1, y_2, \dots, y_n\}.$$

(Discrete)

Example:

$$y = \frac{1}{1 + e^{-(w_0 + w_1 x_1 + w_2 x_2)}}$$



(b)

Classification

- In the **classification**

problem we have

$$C : X \rightarrow Y,$$


$$X \subseteq \mathbb{R}^d,$$

$$Y = \{y_1, y_2, \dots, y_n\}.$$

(Discrete)

Example:

$$y = \frac{1}{1 + e^{-(w_0 + w_1 x_1 + w_2 x_2)}}$$



Clustering

The unsupervised learning problem is similar to the classification one, but the data is not labeled this means that we do not know the category of each example.

- In the **clustering** problem we have

$$C : X \rightarrow T,$$

$$X \subseteq \mathbb{R}^d,$$

$$T = \{t_1, t_2, \dots, t_n\},$$

with unknown n . Example:

$$t = \arg \min_{t_i} \text{distance}(t_i, x)$$

Clustering

The unsupervised learning problem is similar to the classification one, but the data is not labeled this means that we do not know the category of each example.

- In the **clustering** problem we have


$$C : X \rightarrow T,$$

$$X \subseteq \mathbb{R}^d,$$

$$T = \{t_1, t_2, \dots, t_n\},$$

with unknown n . Example:

$$t = \arg \min_{t_i} \text{distance}(t_i, x)$$



Reinforcement Learning

- In reinforcement learning the agent must learn to perform a task. It means going from one start state s_{start} to one final state s_{final} .

Reinforcement Learning


- In reinforcement learning the agent must learn to perform a task. It means going from one start state s_{start} to one final state s_{final} .
- The execution of the task can be seen as a sequence of actions a_1, a_2, \dots, a_n .

Reinforcement Learning

- In reinforcement learning the agent must learn to perform a task. It means going from one start state s_{start} to one final state s_{final} .
- The execution of the task can be seen as a sequence of actions a_1, a_2, \dots, a_n .
- In order to learn the right sequence of actions the agent must interact with its environment selecting at each moment one action a_{t+1} as a function of the current environment state s_t . This is $a_{t+1} = F(s_t)$.


Reinforcement Learning

Markov decision process. $MDP = (S, A, T, R)$



Reinforcement Learning


Markov decision process. $MDP = (S, A, T, R)$



- Set of states
 $S = \{s_1, s_2, \dots\}$

Reinforcement Learning


Markov decision process. $MDP = (S, A, T, R)$



- Set of states
 $S = \{s_1, s_2, \dots\}$
- Set of actions
 $A(s) = \{a_1, a_2, \dots\}$

Reinforcement Learning


Markov decision process. $MDP = (S, A, T, R)$



- Set of states
 $S = \{s_1, s_2, \dots\}$
- Set of actions
 $A(s) = \{a_1, a_2, \dots\}$
- Transition function
(dynamics of the environment)
 $T : S \times A \rightarrow S$

Reinforcement Learning

Markov decision process. $MDP = (S, A, T, R)$





- Set of states
 $S = \{s_1, s_2, \dots\}$
- Set of actions
 $A(s) = \{a_1, a_2, \dots\}$
- Transition function
(dynamics of the environment)
 $T : S \times A \rightarrow S$
- Reward function
 $R : S \rightarrow \mathbb{R}$

Possible RL solution to the maze problem

We discretize states and actions: 64 states and 4 actions.

Possible RL solution to the maze problem

We discretize states and actions: 64 states and 4 actions.



Reward Function:

$r = -1$ if state is not terminal
 $r = 10$ if state is terminal

Function $V(s)$

S	$V(s)$
1	0
2	0
3	0
4	0
...	...
64	0

Function $V(s)$

S	$V(s)$
1	0
2	0
3	0
4	0
...	...
64	0

S	$V(s)$
1	-20
2	-20
3	-5
4	-22
...	...
64	-24

Function $Q(s,a)$

S	A	$Q(s, a)$
1	1	0
1	2	0
1	3	0
1	4	0
2	1	0
2	2	0
2	3	0
...
64	4	0

Function $Q(s,a)$

S	A	$Q(s, a)$
1	1	0
1	2	0
1	3	0
1	4	0
2	1	0
2	2	0
2	3	0
...
64	4	0

S	A	$Q(s, a)$
1	1	-20
1	2	-23
1	3	-5
1	4	-22
2	1	-20
2	2	-5
2	3	-22
...
64	4	-24

Q-learning Algorithm

Q-learning Algorithm

Initialize $Q(s, a)$ arbitrarily

Q-learning Algorithm

Initialize $Q(s, a)$ arbitrarily
Repeat (for each episode):

Q-learning Algorithm

Initialize $Q(s, a)$ arbitrarily
Repeat (for each episode):
 Initialize s

Q-learning Algorithm

Initialize $Q(s, a)$ arbitrarily

Repeat (for each episode):

 Initialize s

 Repeat (for each step of episode):

Q-learning Algorithm

Initialize $Q(s, a)$ arbitrarily

Repeat (for each episode):

 Initialize s

 Repeat (for each step of episode):

 Choose a from s using policy derived from Q (ϵ -greedy)

Q-learning Algorithm

Initialize $Q(s, a)$ arbitrarily

Repeat (for each episode):

 Initialize s

 Repeat (for each step of episode):

 Choose a from s using policy derived from Q (ϵ -greedy)

 Take action a , observe r, s'

Q-learning Algorithm

Initialize $Q(s, a)$ arbitrarily

Repeat (for each episode):

Initialize s

Repeat (for each step of episode):

Choose a from s using policy derived from Q (ϵ -greedy)

Take action a , observe r, s'

$$Q(s, a) \leftarrow Q(s, a) + \alpha [r + \gamma \max_{a'} Q(s', a') - Q(s, a)]$$

Q-learning Algorithm

Initialize $Q(s, a)$ arbitrarily

Repeat (for each episode):

 Initialize s

 Repeat (for each step of episode):

 Choose a from s using policy derived from Q (ϵ -greedy)

 Take action a , observe r, s'

$Q(s, a) \leftarrow Q(s, a) + \alpha[r + \gamma \max_{a'} Q(s', a') - Q(s, a)]$

$s \leftarrow s'$

Q-learning Algorithm

Initialize $Q(s, a)$ arbitrarily

Repeat (for each episode):

 Initialize s

 Repeat (for each step of episode):

 Choose a from s using policy derived from Q (ϵ -greedy)

 Take action a , observe r, s'

$Q(s, a) \leftarrow Q(s, a) + \alpha[r + \gamma \max_{a'} Q(s', a') - Q(s, a)]$

$s \leftarrow s'$

 Until s is terminal

Some videos

- Inverted pendulum
- Object handling

Thank you!

Dr. Víctor Uc Cetina
cetina@informatik.uni-hamburg.de